

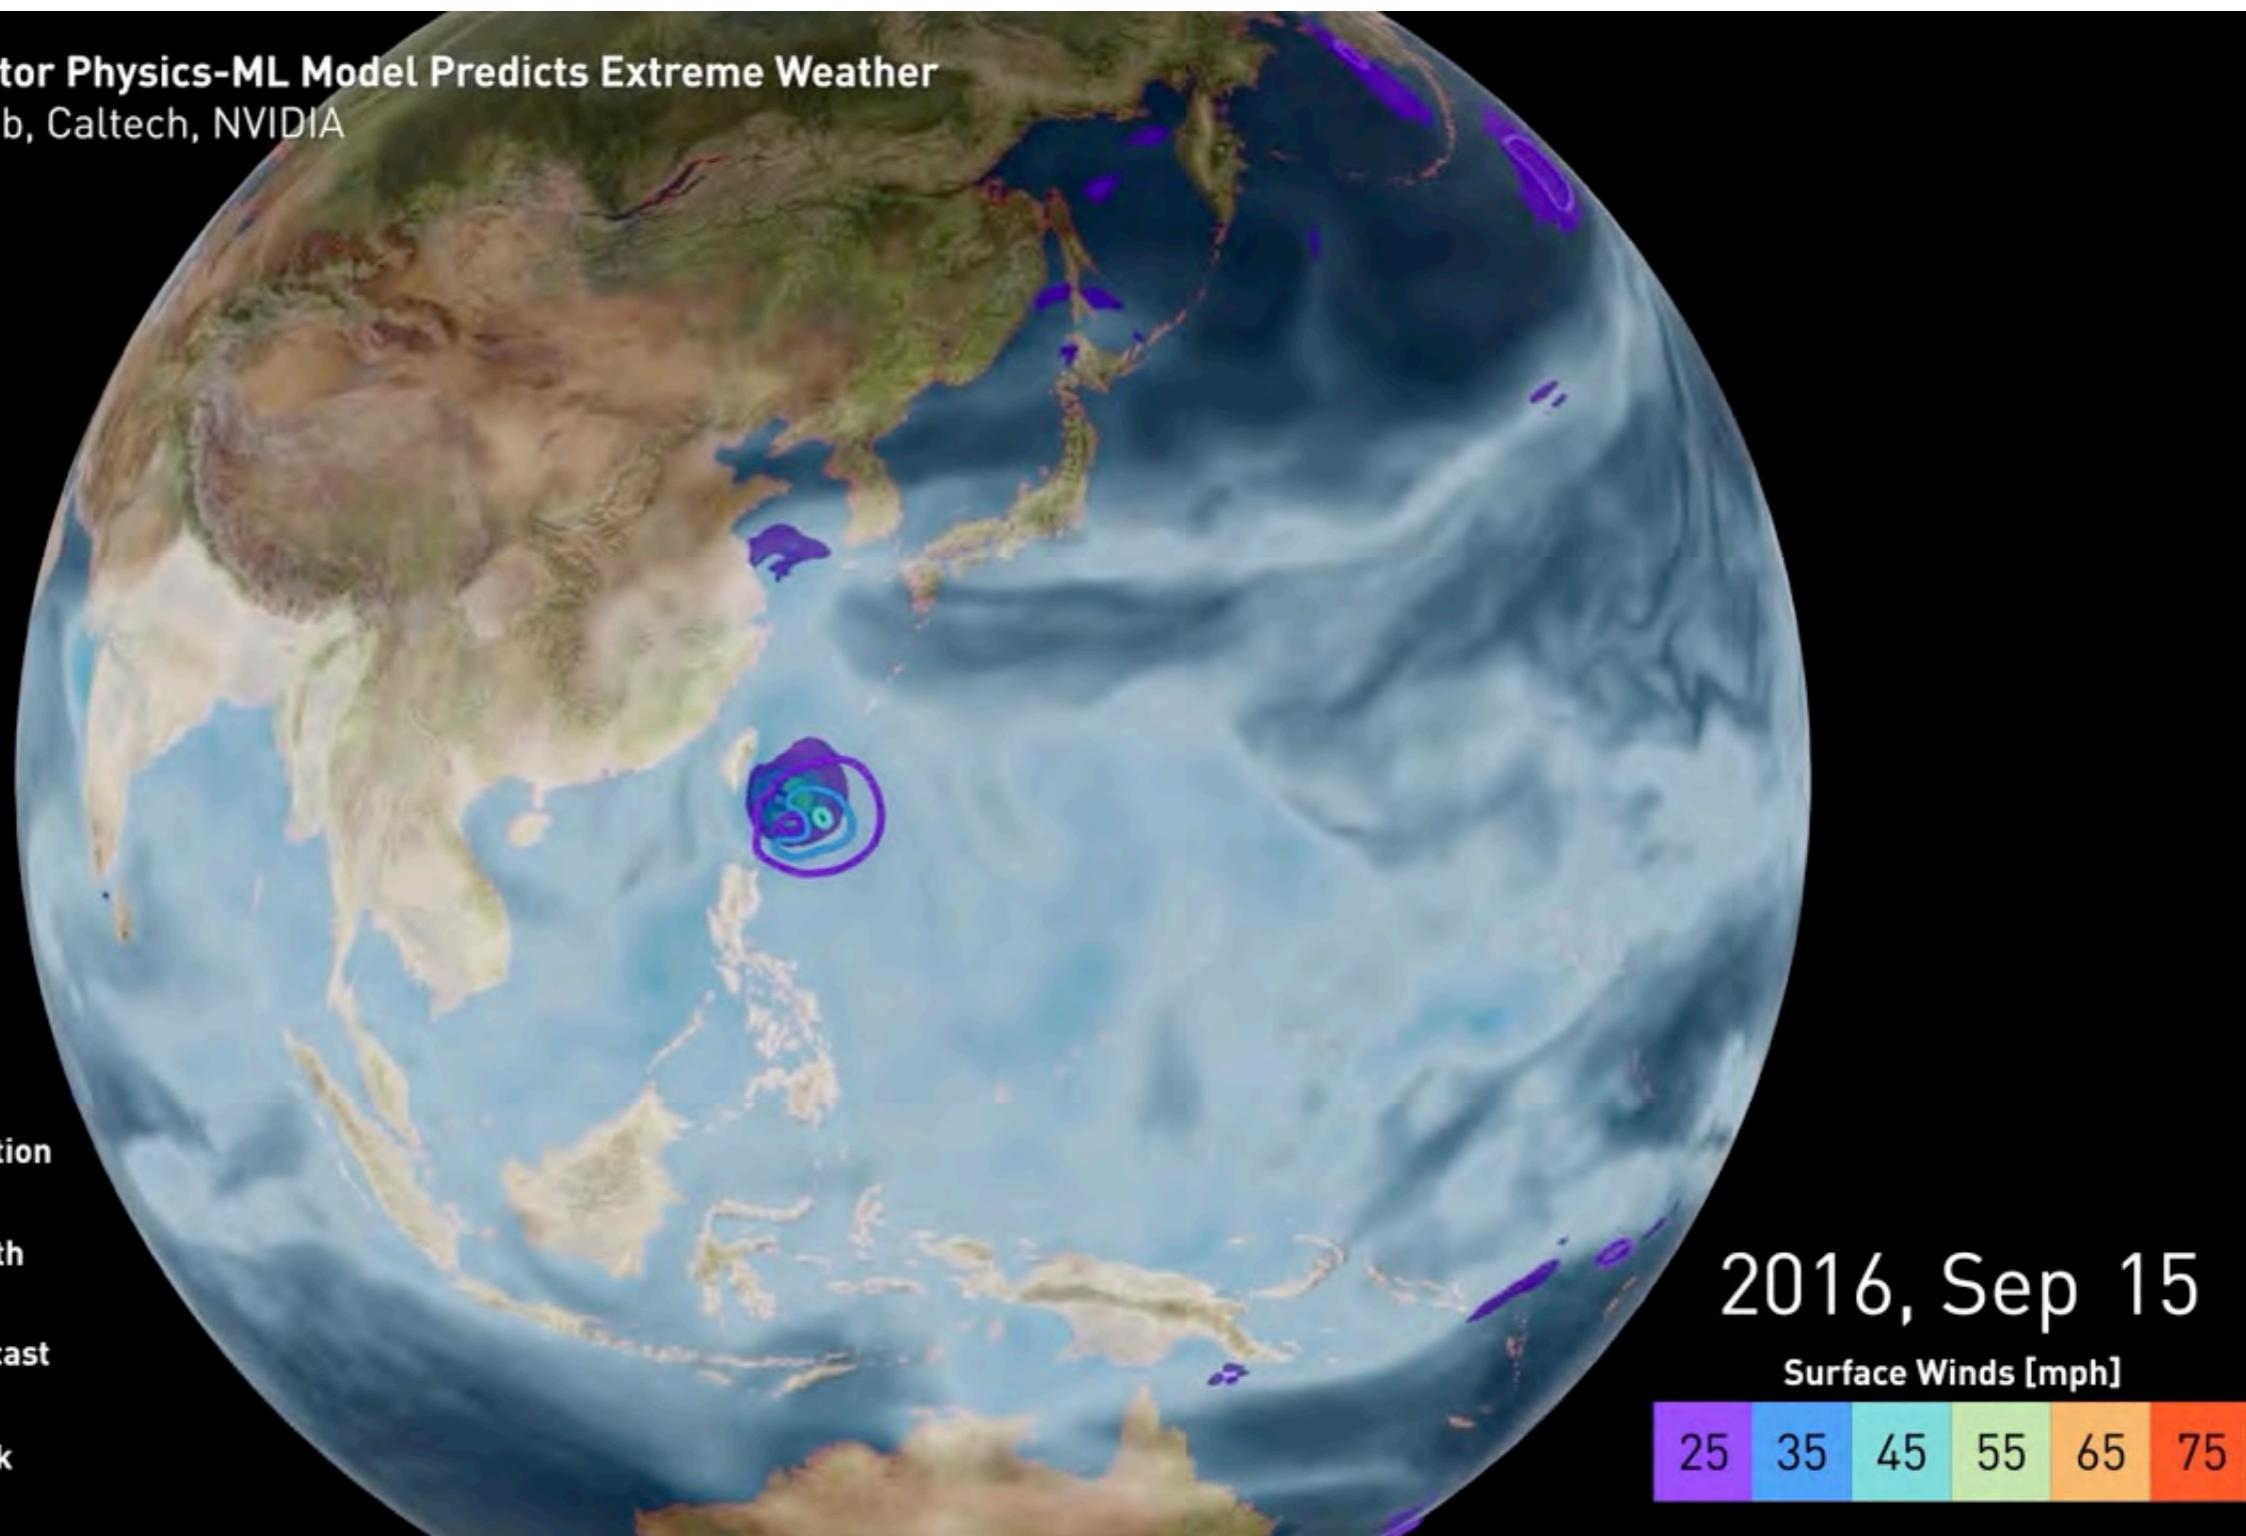
# MEAM4600:AI for Science and Engineering

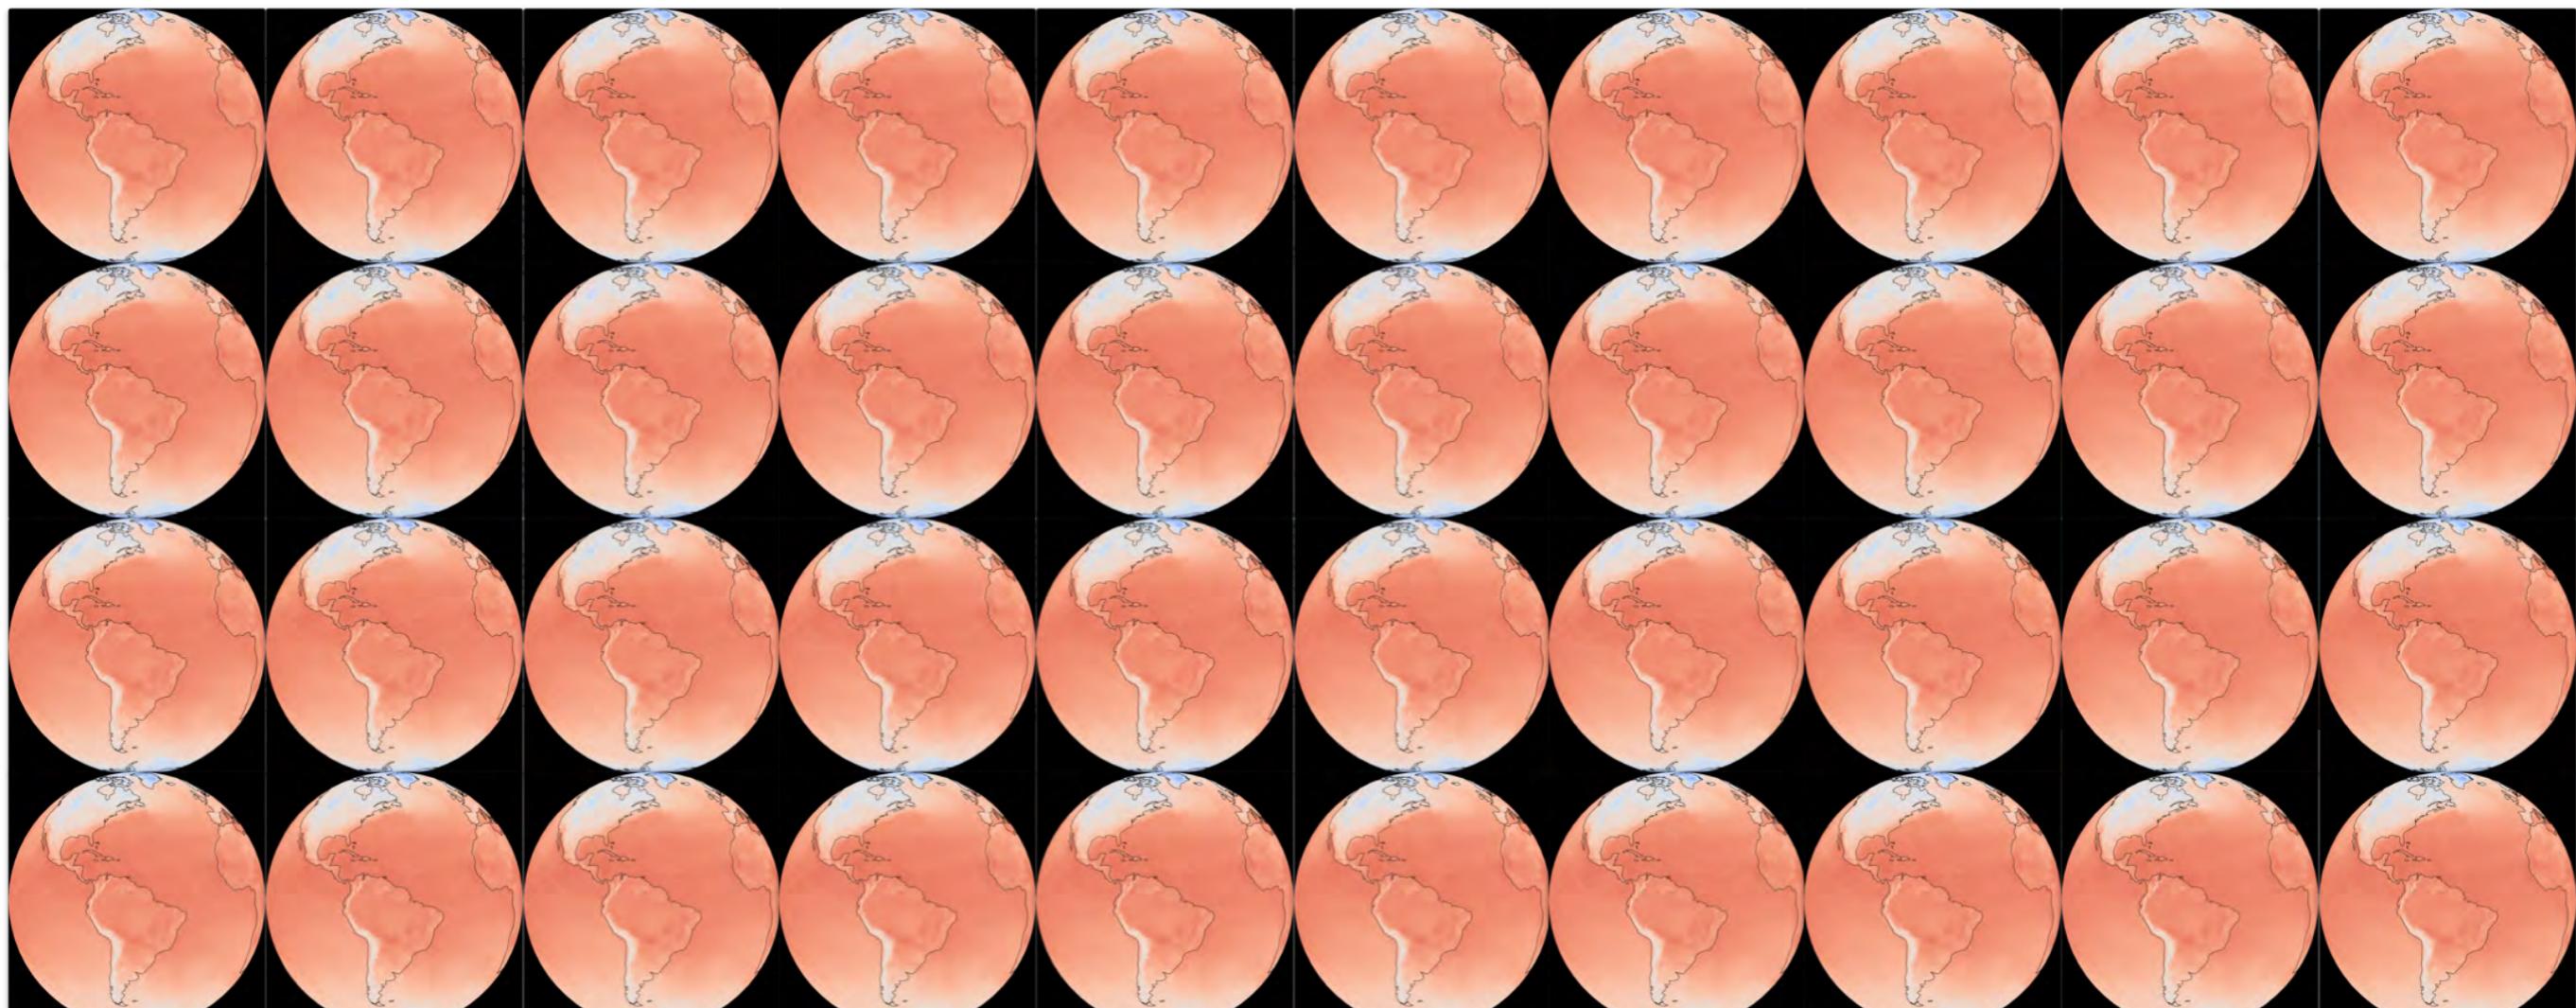
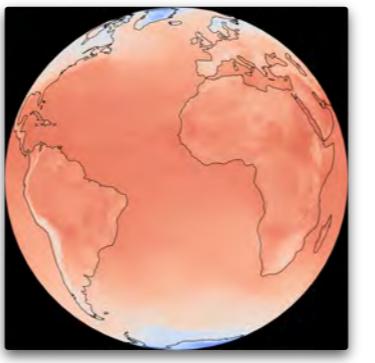
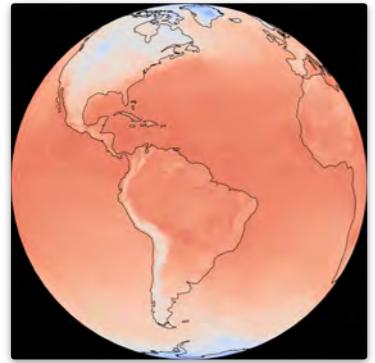
*Lecture #0: Introduction, motivation and course logistics*



# Predictive science and engineering

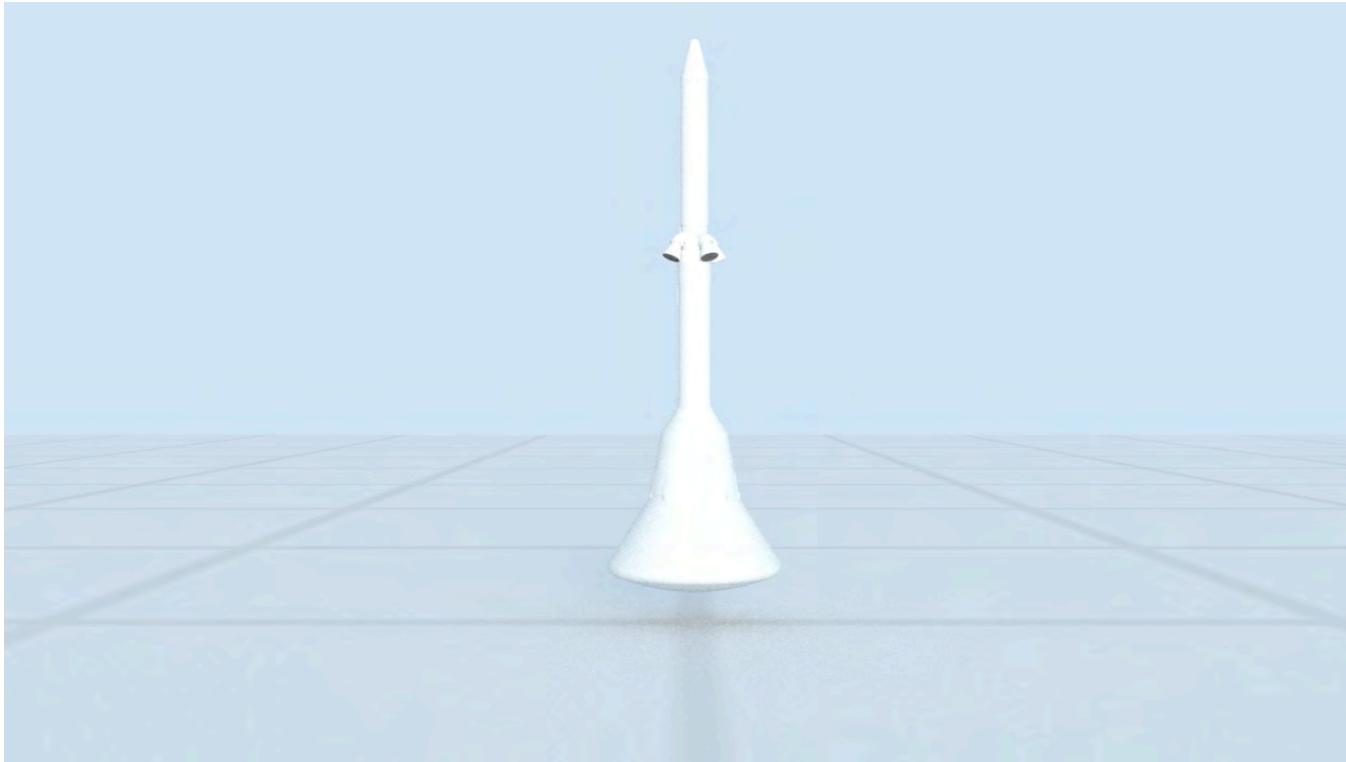
Fourier Neural Operator Physics-ML Model Predicts Extreme Weather  
Lawrence Berkeley Lab, Caltech, NVIDIA



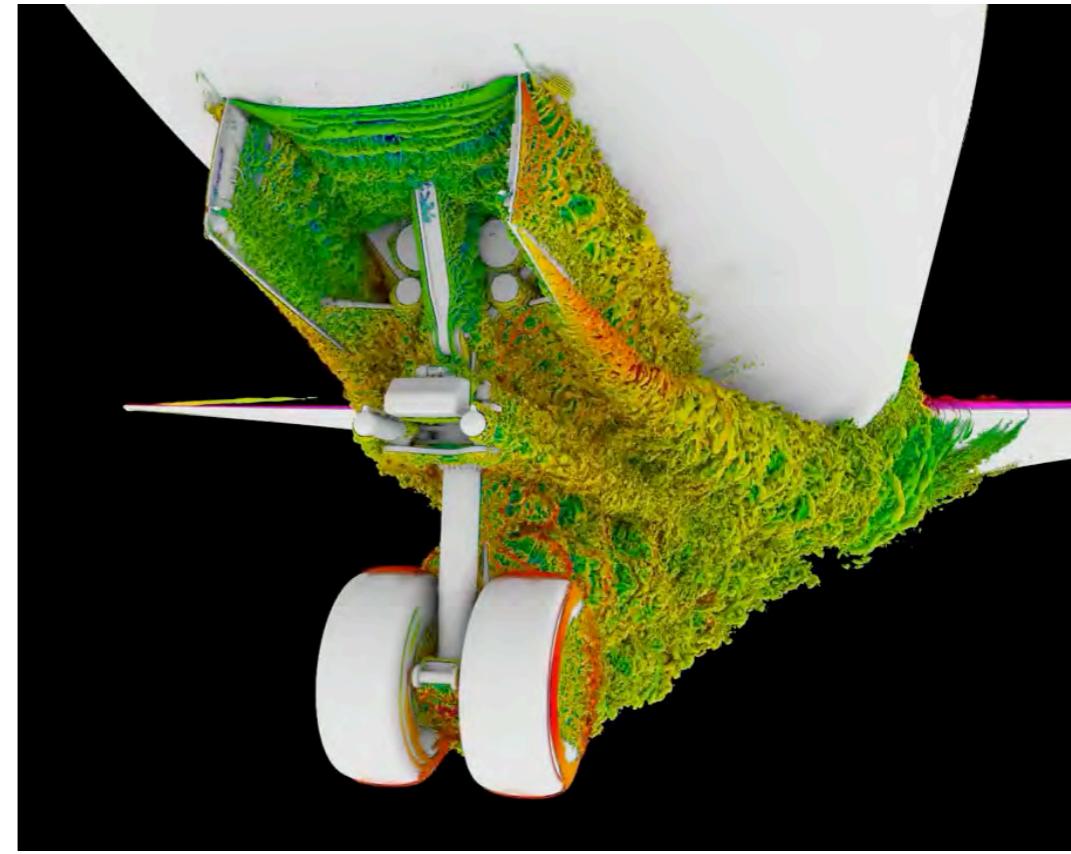


# Computational science & engineering

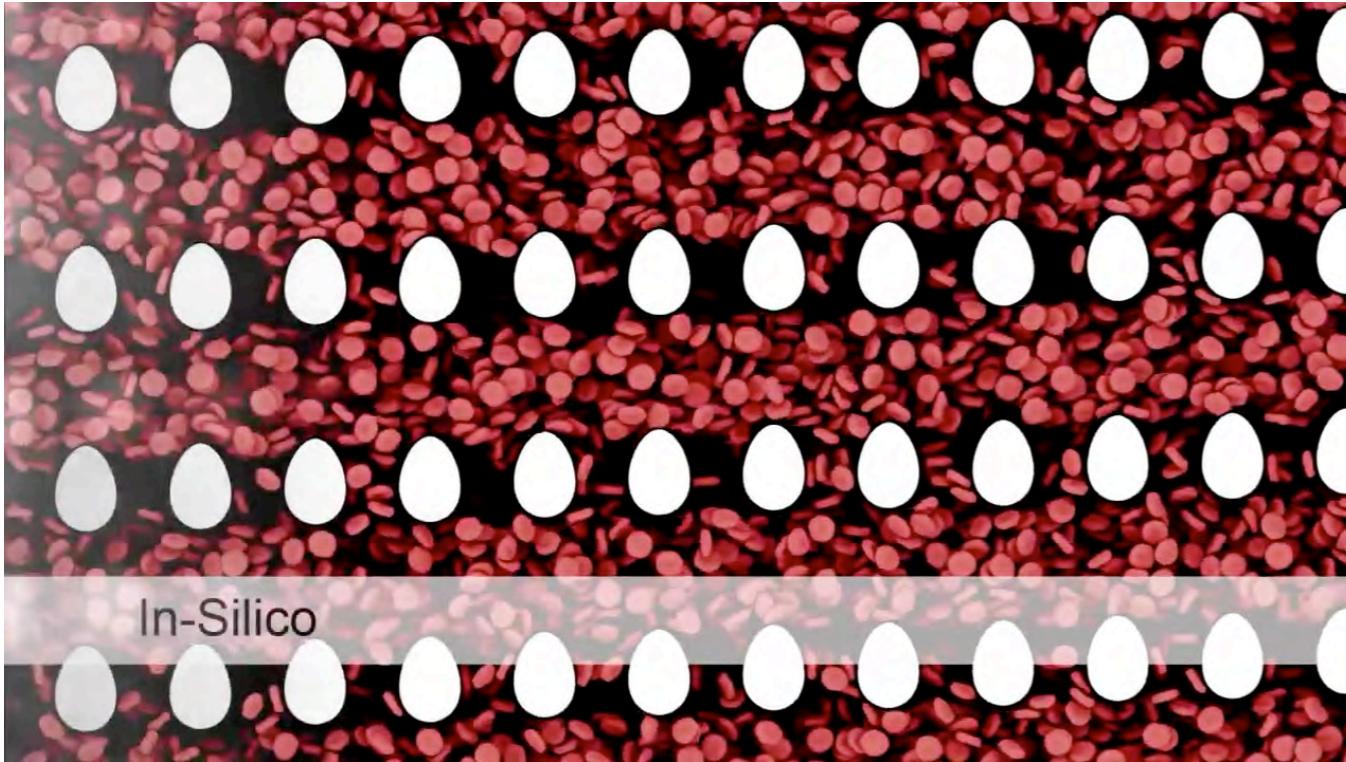
*Exploiting the power of computation to resolve major challenges at the frontiers of engineering, natural and life sciences*



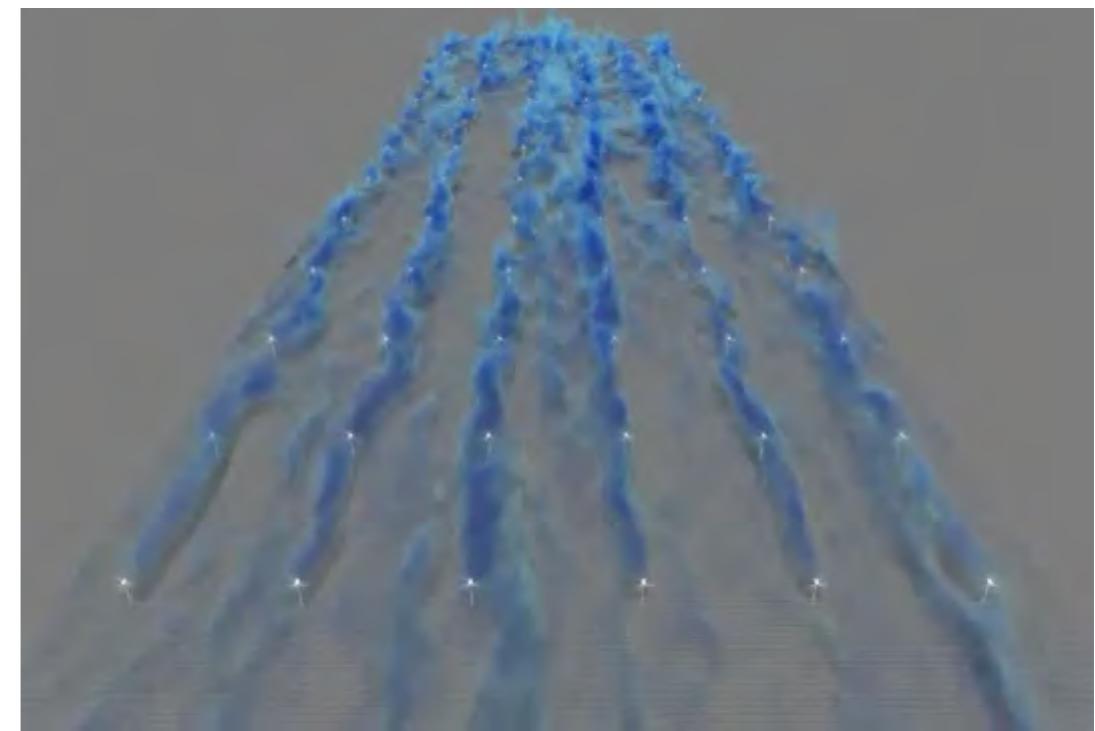
Simulation of Orion Spacecraft Launch Abort System (NASA Ames)



Detailed flow around an aircraft's landing gear (NASA Ames)



In-Silico Lab-on-a-Chip: high-throughput simulations of micro-fluidics at cell resolution (ETH Zurich)



Flow field in a simulated wind-farm (JHU)

# Roadmap to Exascale computing

	2012	2016	2020	2024
Peak flops	10-20 PF	100-200 PF	500-2000 PF	2000 - 4000 PF
Memory	0.5-1 PB	5-10 PB	32 – 64 PB	50-100 PB
Burst storage bandwidth	NA	5 TB/s	32 TB/s	50 TB/s
Burst capacity (cache)	NA	500 TB	3 PB	5 PB
Mid-tier capacity (disk)	20 PB	100 PB	1 EB	5 EB
Bottom-tier capacity (tape)	100 PB	1 EB	10 EB	50 EB
I/O servers	400	500	600	700



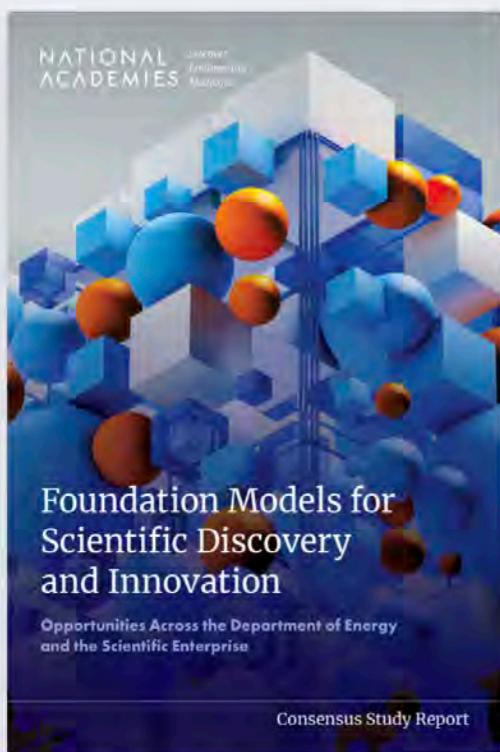
| ExaFlop:  $10^{18}$  floating point operations per second.

# Roadmap to Trillion Sensors



Sensors are expected to generate a bronto-bytes (1,000 trillion trillion) of data by 2025!

# An Imminent Paradigm Shift



CONSENSUS

Foundation Models for Scientific Discovery and Innovation:  
Opportunities Across the Department of Energy and the  
Scientific Enterprise

2025

Read Online

Download PDF



Paperback - \$21.60

Buy the book

Request an exam or desk copy



Foundation models – artificial intelligence systems trained on massive data sets to perform a wide range of tasks – have the potential to transform scientific discovery and innovation. At the request of the U.S. Department of Energy (DOE), the National Academies conducted a study to consider the capabilities of current foundation models as well as future possibilities and challenges. Foundation Models for Scientific Discovery and Innovation explores how foundation models can complement traditional computational methods to advance scientific discovery, highlights successful use cases, and recommends strategic approaches and investments to support DOE's mission.

# An Imminent Paradigm Shift

The US just launched a \$100 Billion Manhattan Project for AI called Genesis

## THE GENESIS MISSION

The Largest Scientific Reorganization Since Apollo



November 24, 2025

Doubling American Scientific Productivity in 10 Years

### THE EXASCALE ARSENAL (AI Focus)



### THE INFRASTRUCTURE MAP (National Scale)



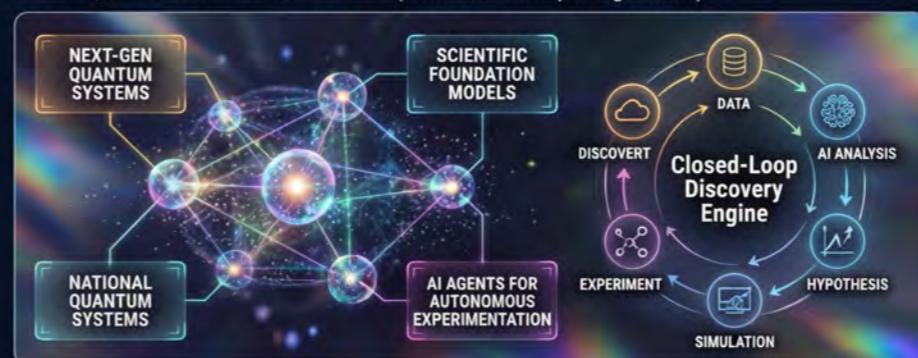
### HISTORICAL COMPARISON TIMELINE

Manhattan Project	Apollo Program	Genesis Mission
1942-1945 Cost: \$30B (adjusted) Duration: 3 Years Workforce: 130,000 Output: Atomic Bomb	1961-1972 Cost: \$257B (adjusted) Duration: 12 Years Workforce: 400,000 Output: Moon Landing	2025-2035 Cost: \$100B+ Duration: 10 Years Workforce: 40,000+ Output: AI Science Platform

### THE ENERGY CHALLENGE (Energy Focus)



### THE QUANTUM FRONTIER (Quantum Computing Focus)



### MISSION OBJECTIVES (Three Pillars)



### COUNTDOWN TIMELINE



Sources: DOE, White House, TOP500, AWS

THE 21ST CENTURY WILL BE DECIDED BY COMPUTATIONAL SUPREMACY.

# Autonomy



Tesla AI Day: [https://www.youtube.com/watch?v=ODSJsviD\\_SU](https://www.youtube.com/watch?v=ODSJsviD_SU)

# Intelligent experimentation

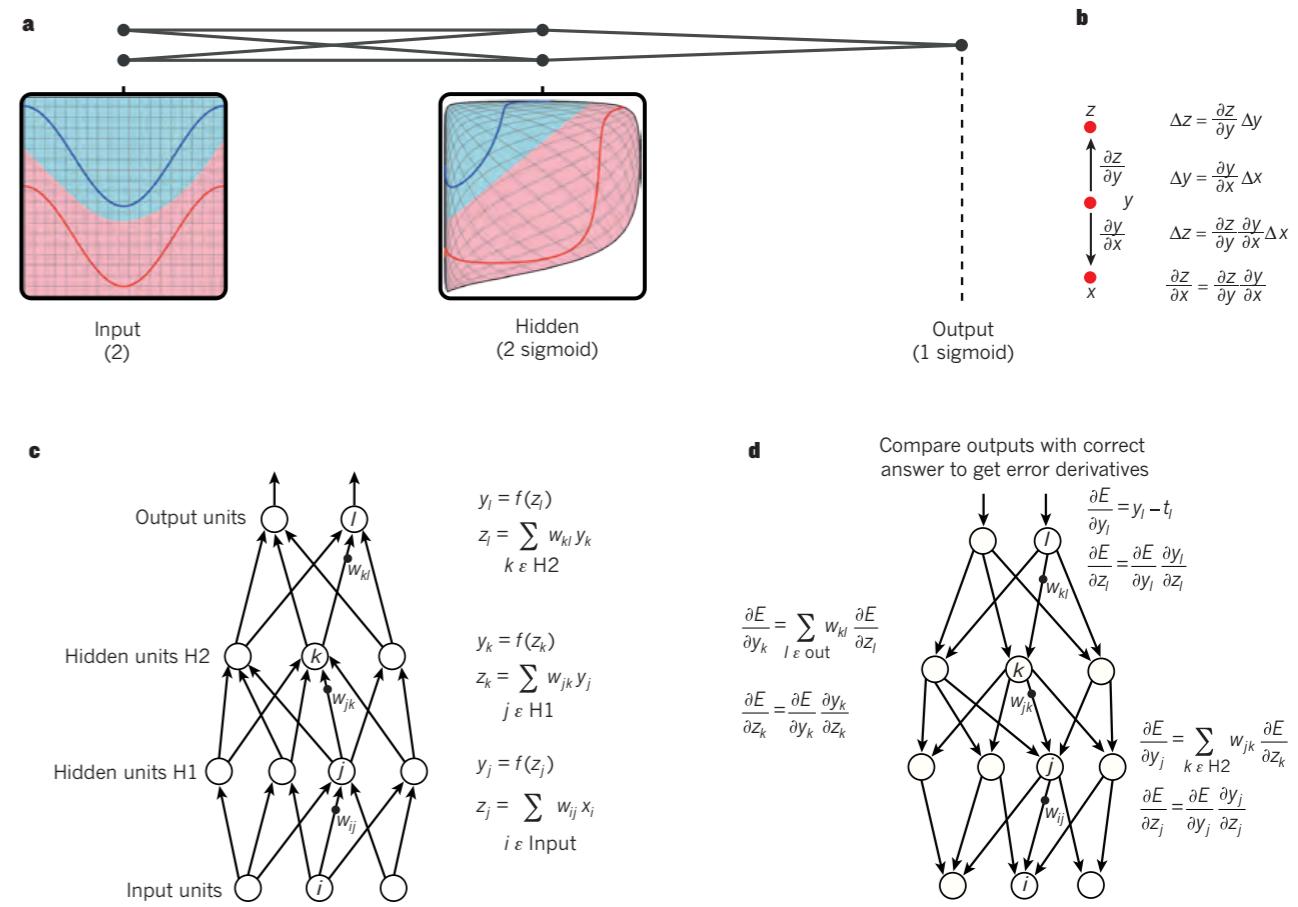
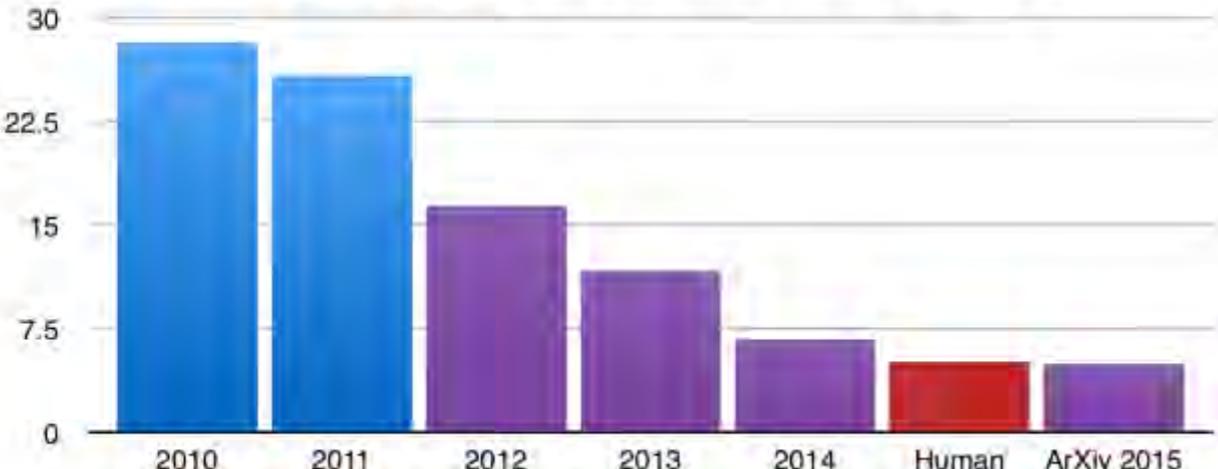


*“In its first year of operation, the Intelligent Towing Tank (ITT) conducted about 100,000 total experiments, essentially completing the equivalent of a PhD student’s five years’ worth of experiments in a matter of weeks.”*

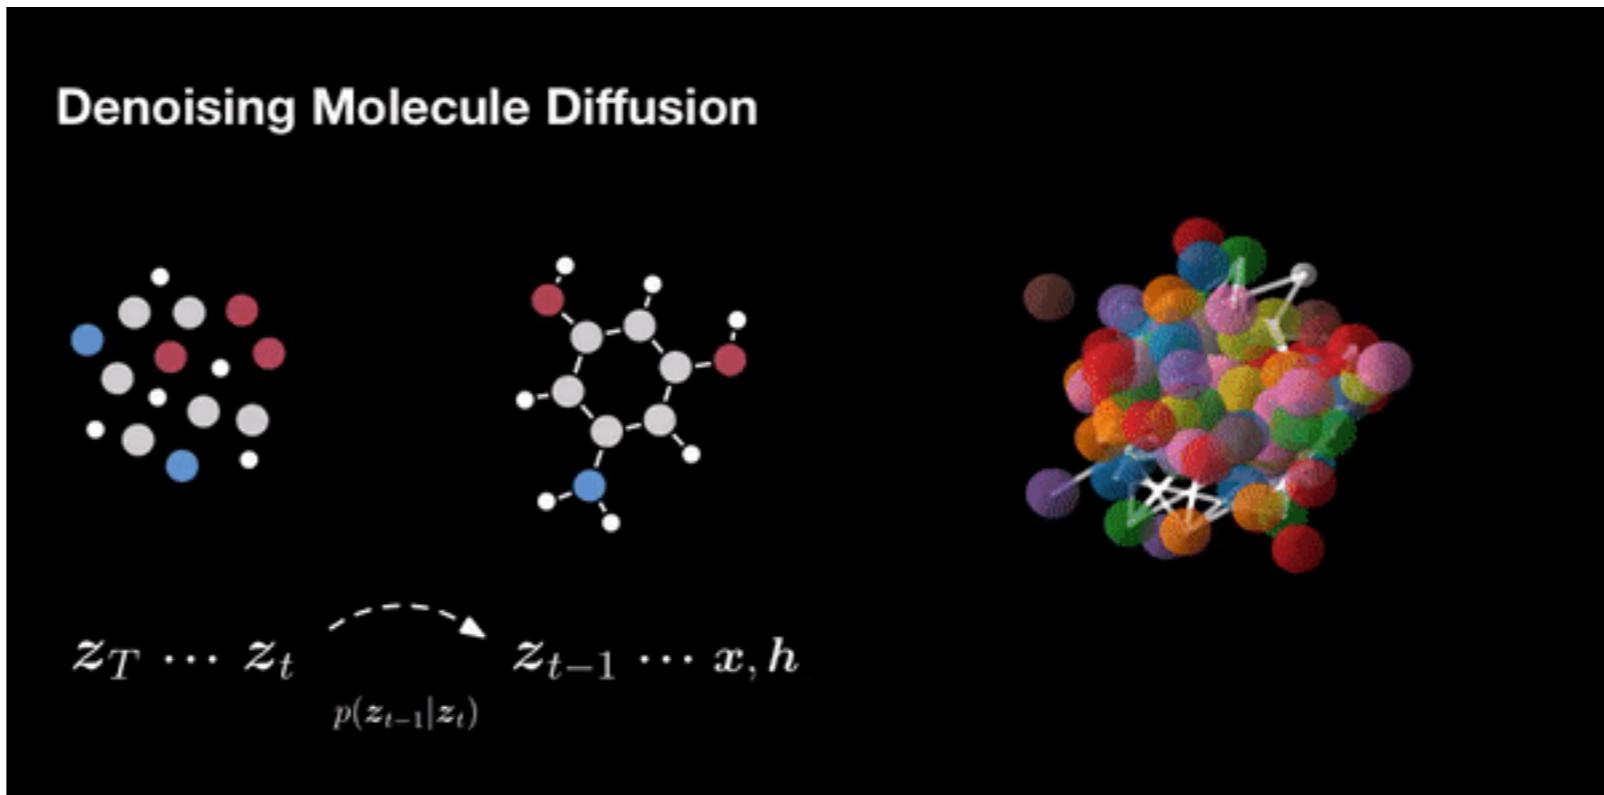
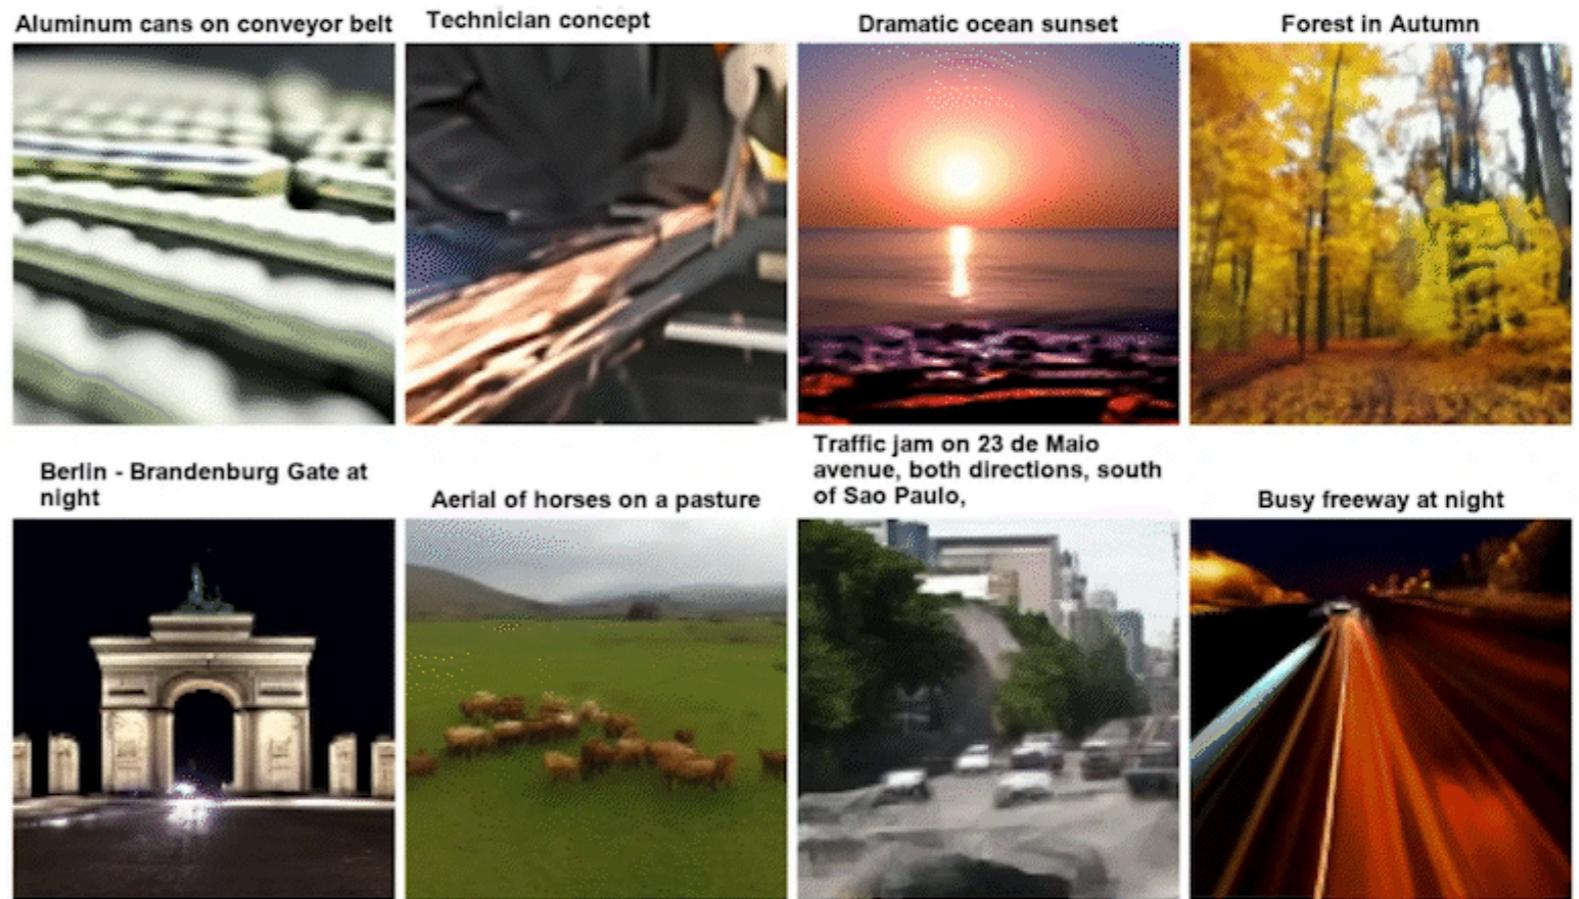
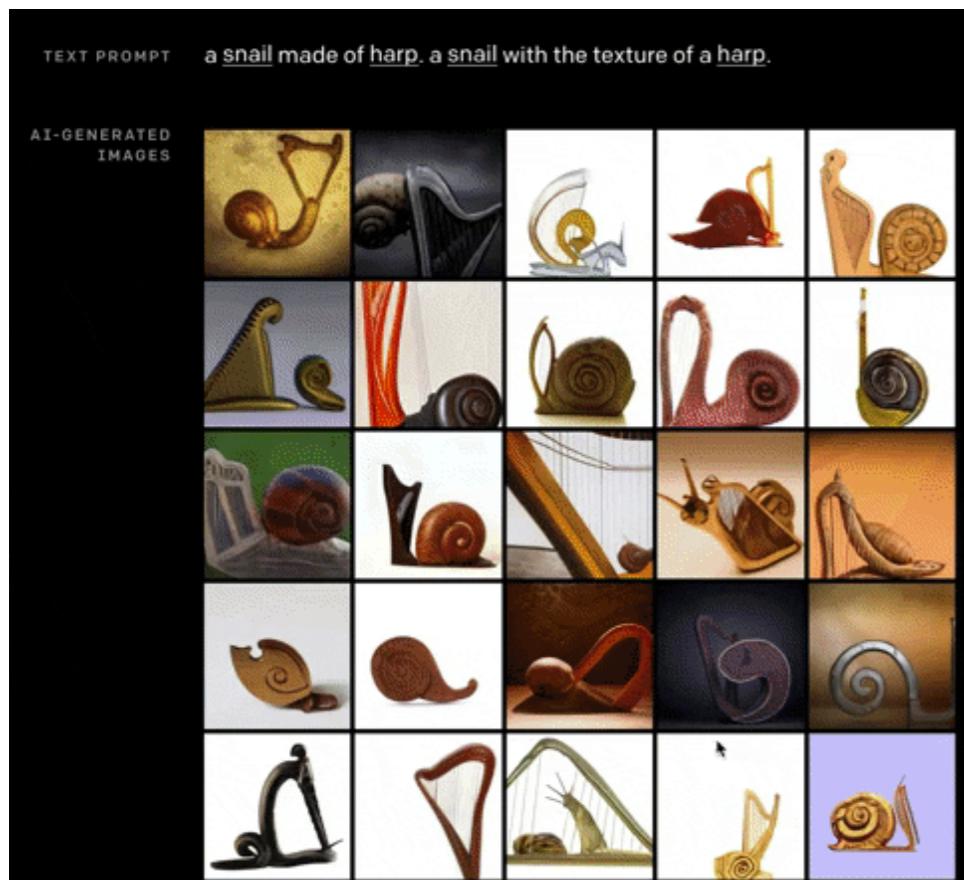
<https://news.mit.edu/2019/intelligent-towing-tank-propels-research-1209>

# Recent success of machine learning

ILSVRC top-5 error on ImageNet



# Working with high-dimensional data



# Working with multiple data modalities

Gemini

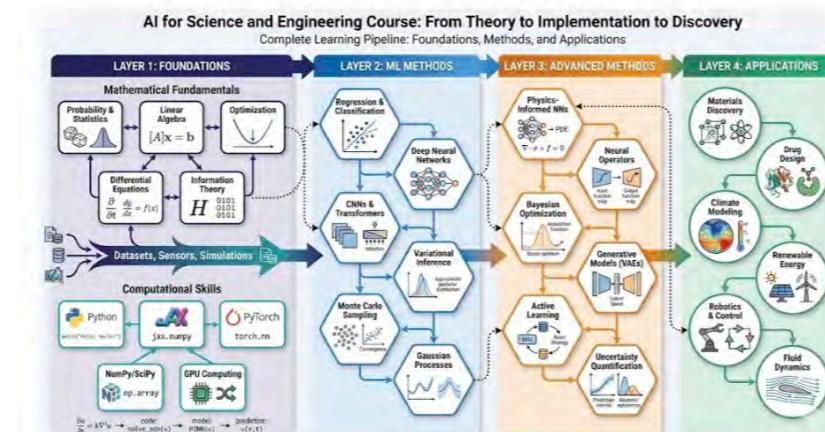
Create a clean, modern academic infographic illustrating an AI for Science a...



Create a clean, modern academic infographic illustrating an AI for Science and Engineering course with complete learning pipeline.

LAYOUT: Four-layer hierarchical flow diagram (left to...

Show thinking (Nano Banana Pro)



Like Dislike Comment More

Describe your image

+ - Image X

Thinking

# Data-driven science & engineering

Neural gigapixel images



Neural SDF



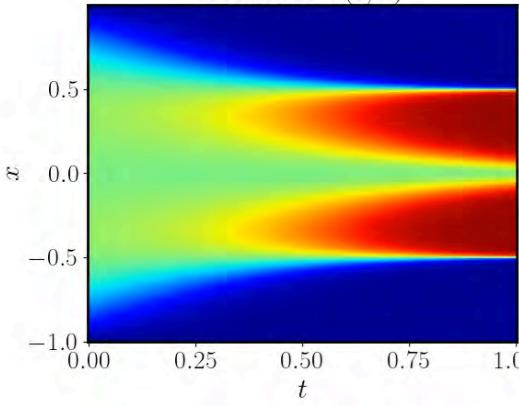
NeRF



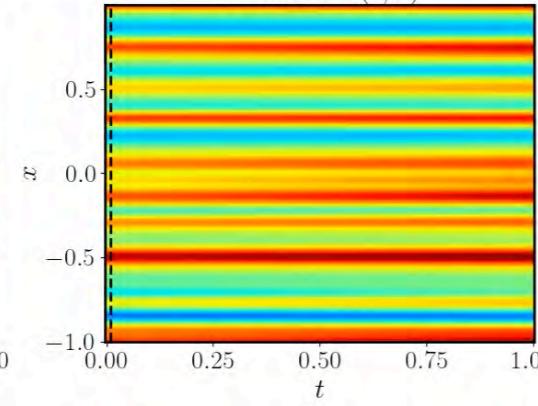
Neural volume

Elapsed training time: 0 seconds

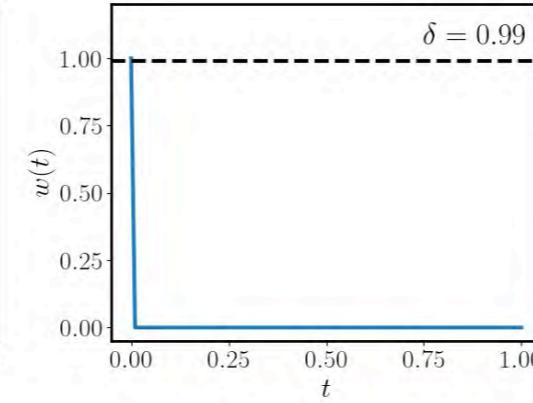
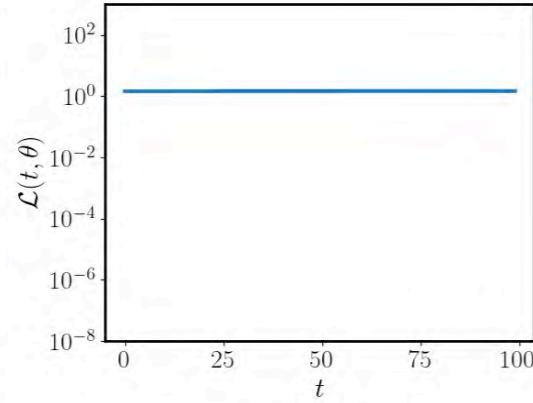
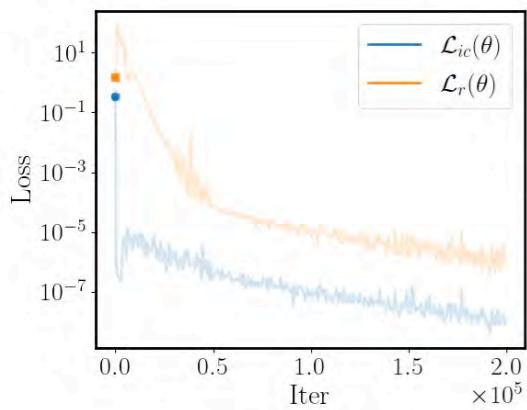
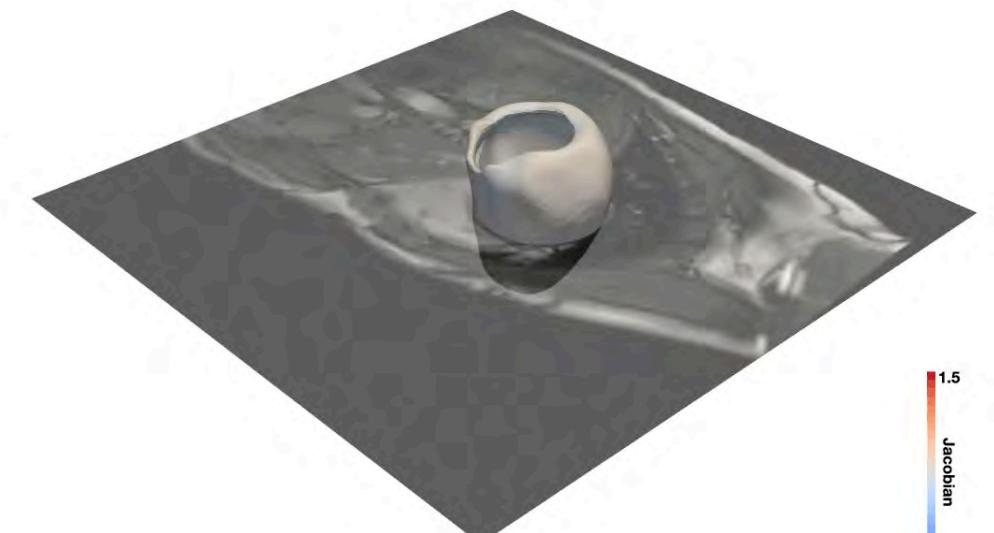
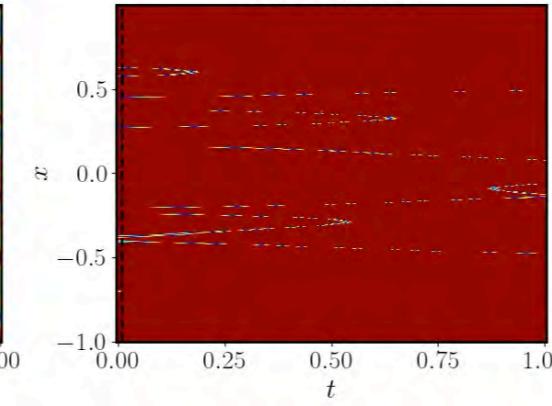
Reference  $u(t, x)$



Predicted  $u(t, x)$

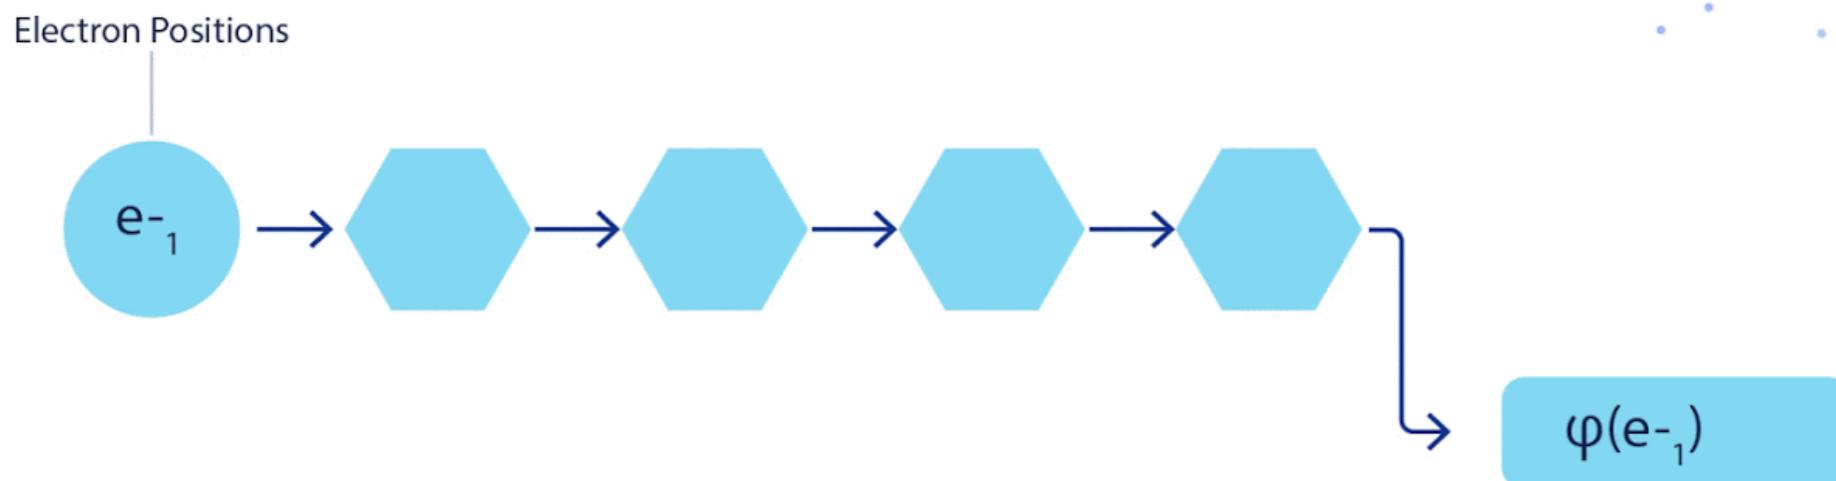
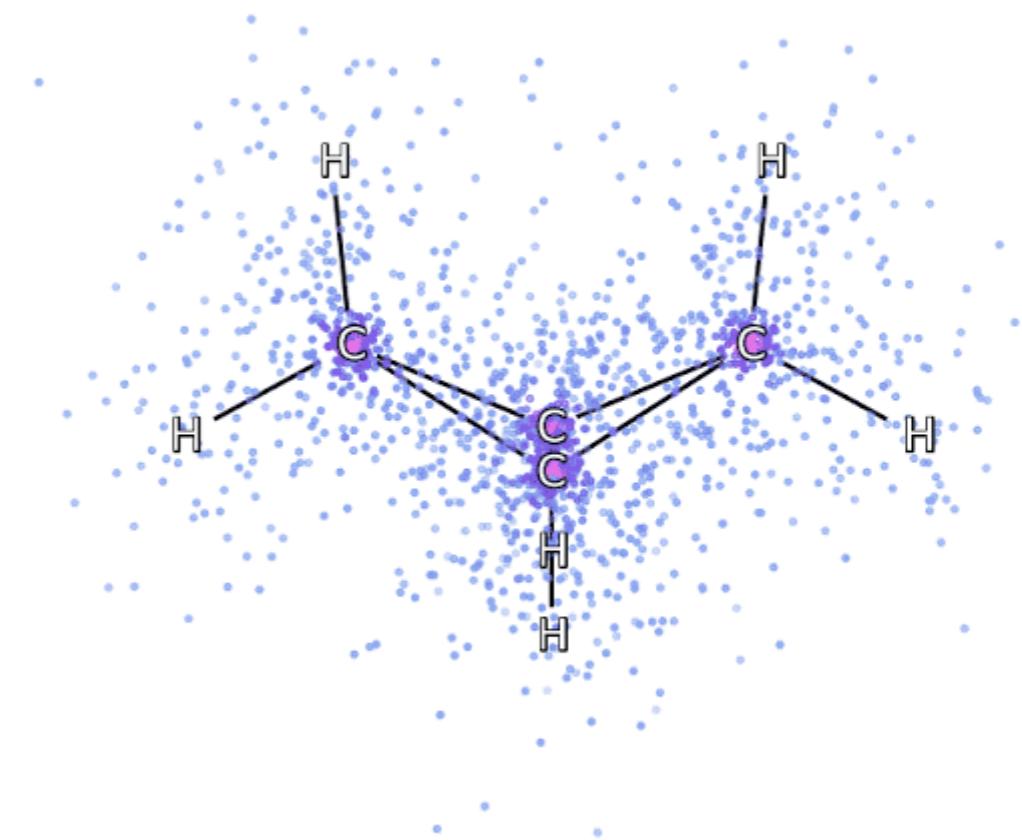
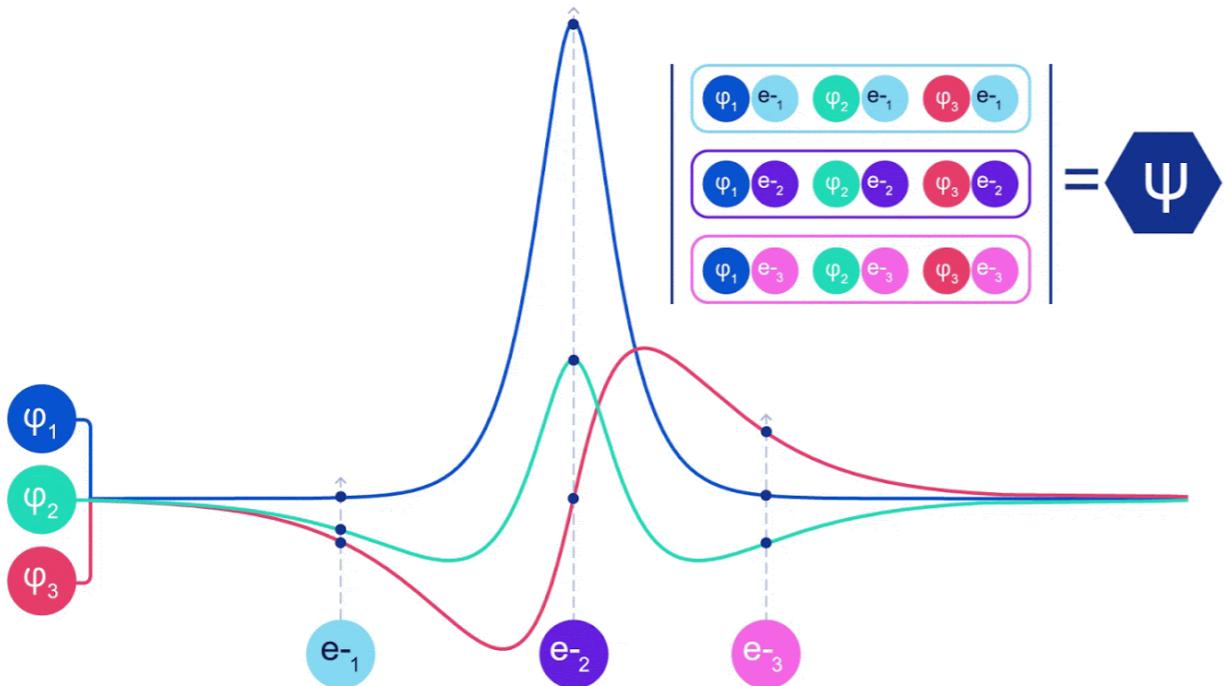


Absolute error

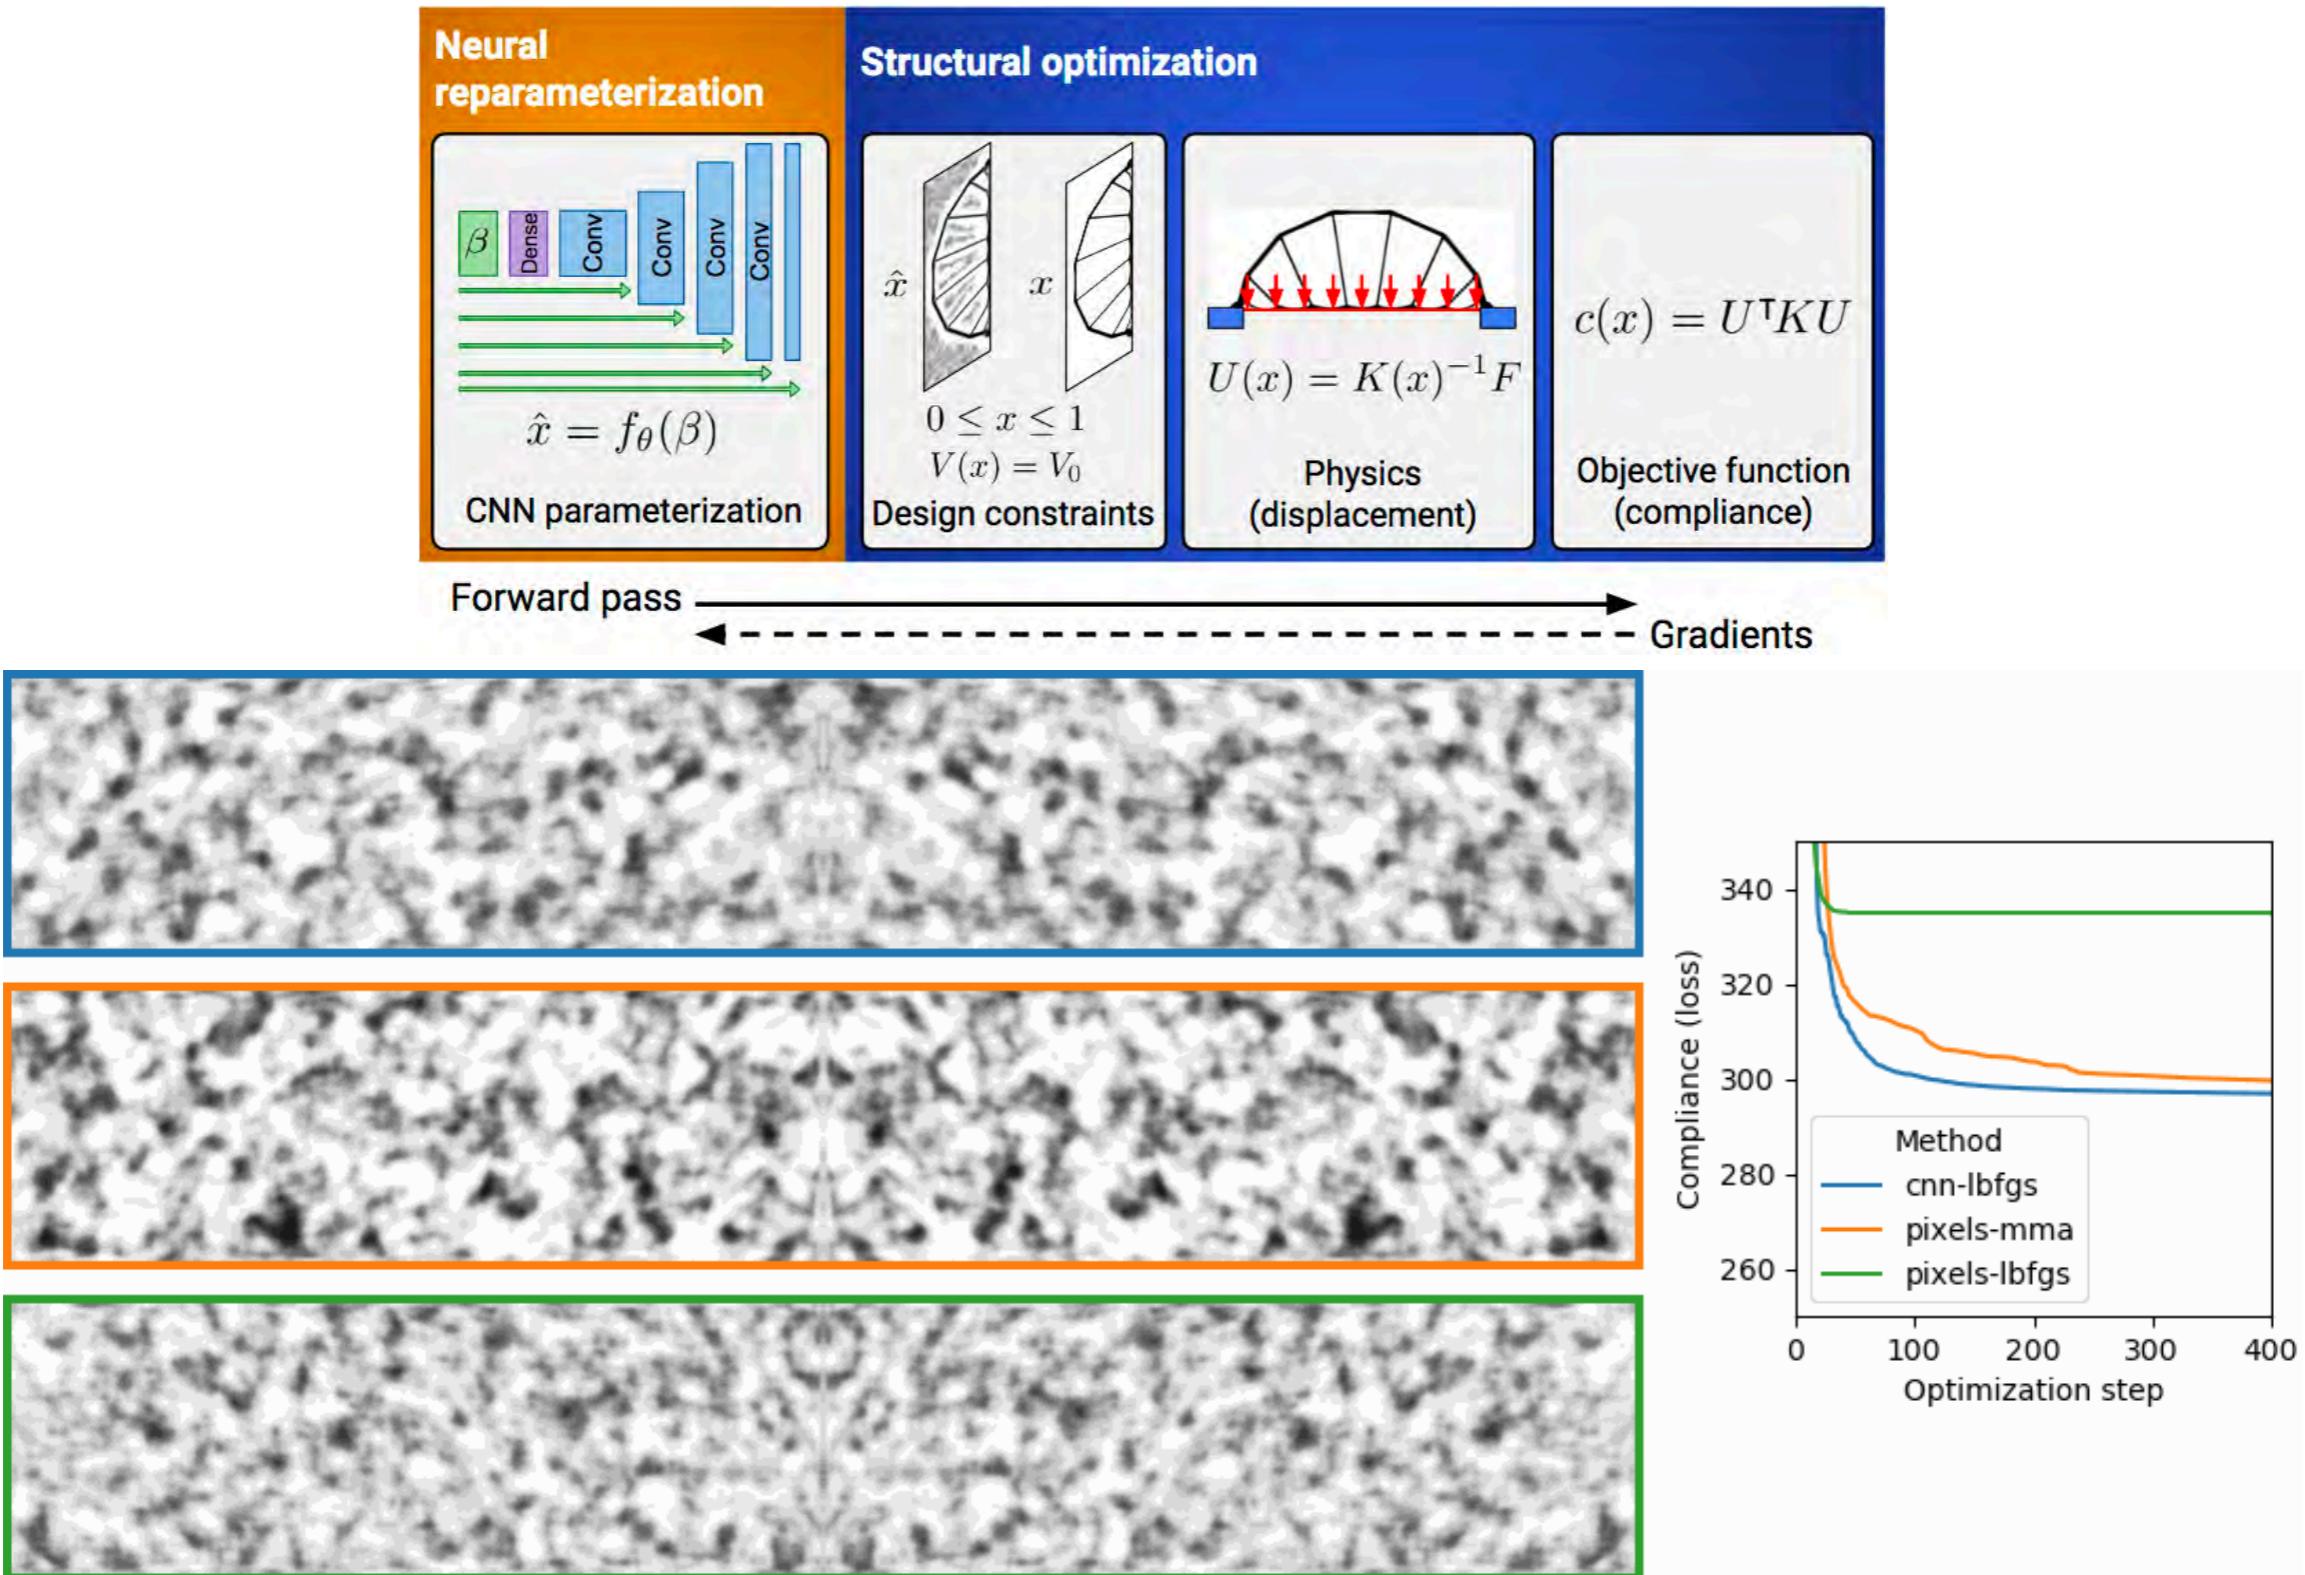


Jacobian  
1.5  
0.5

# Data-driven science & engineering



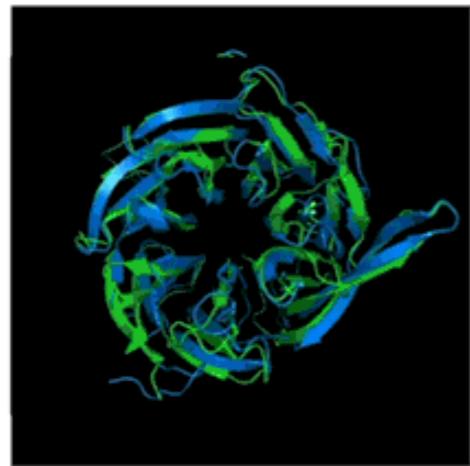
# Data-driven science & engineering



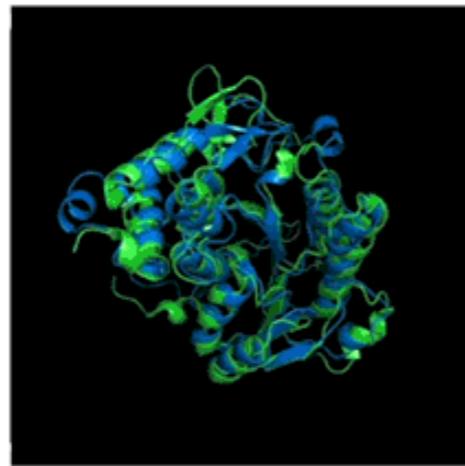
Optimizing a bridge structure. In the top frame, optimization happens in the weight space of a CNN. In the next two frames it happens on a finite element grid.

# Data-driven science & engineering

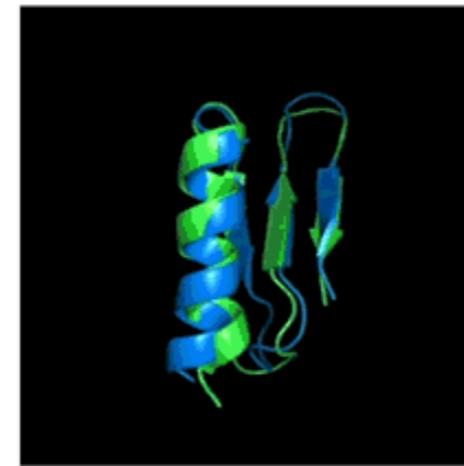
T0954 / 6CVZ



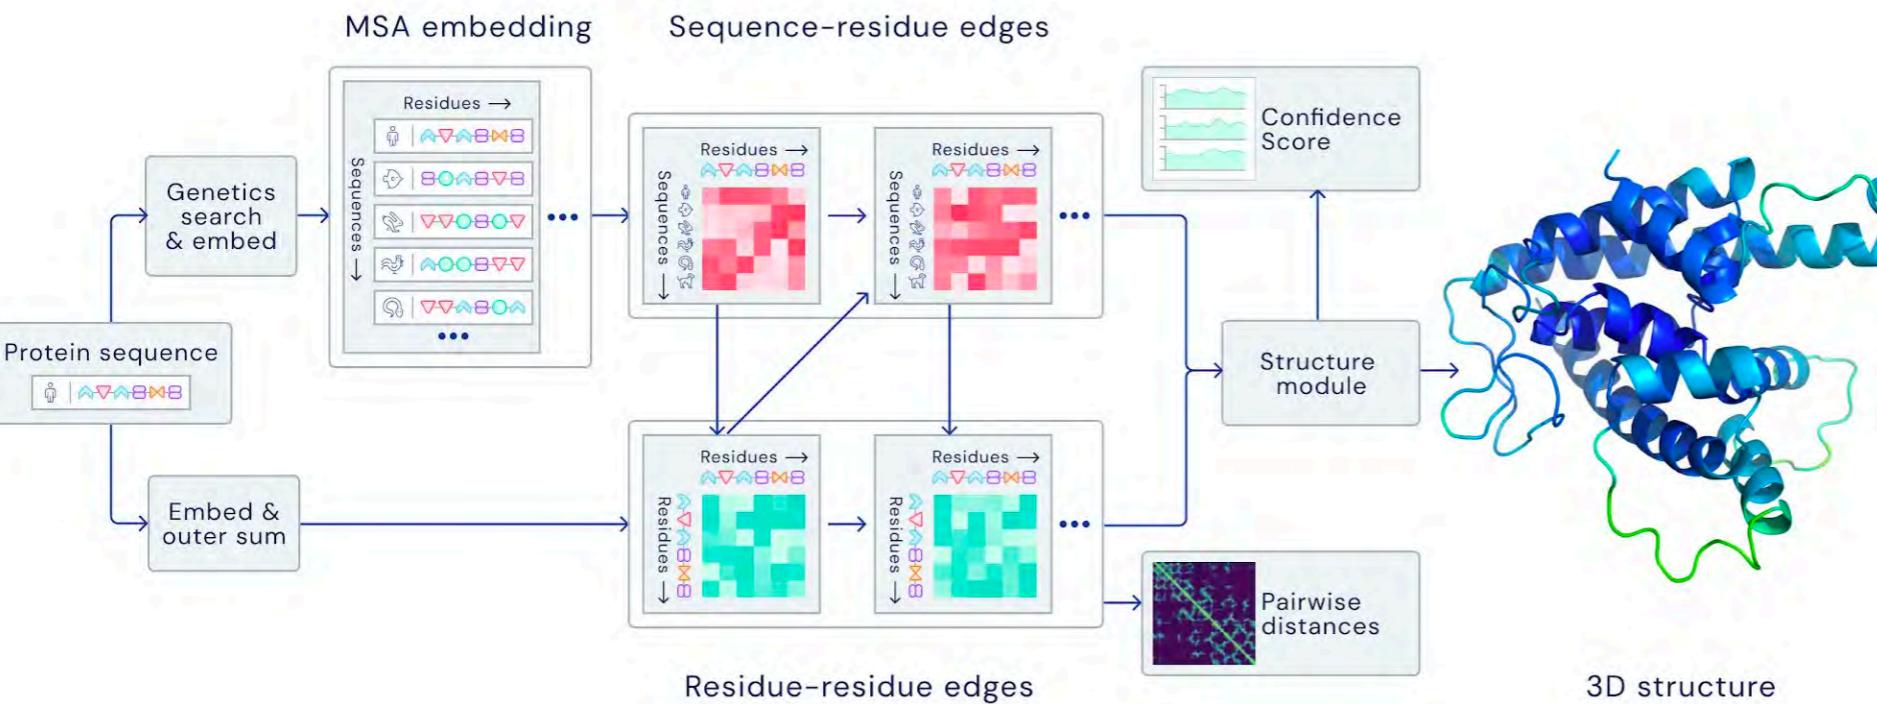
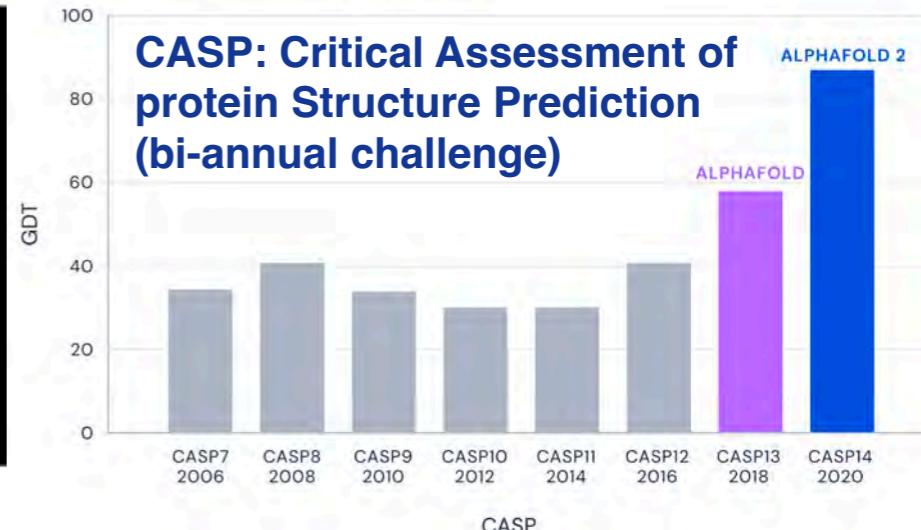
T0965 / 6D2V



T0955 / 5W9F



Median Free-Modelling Accuracy



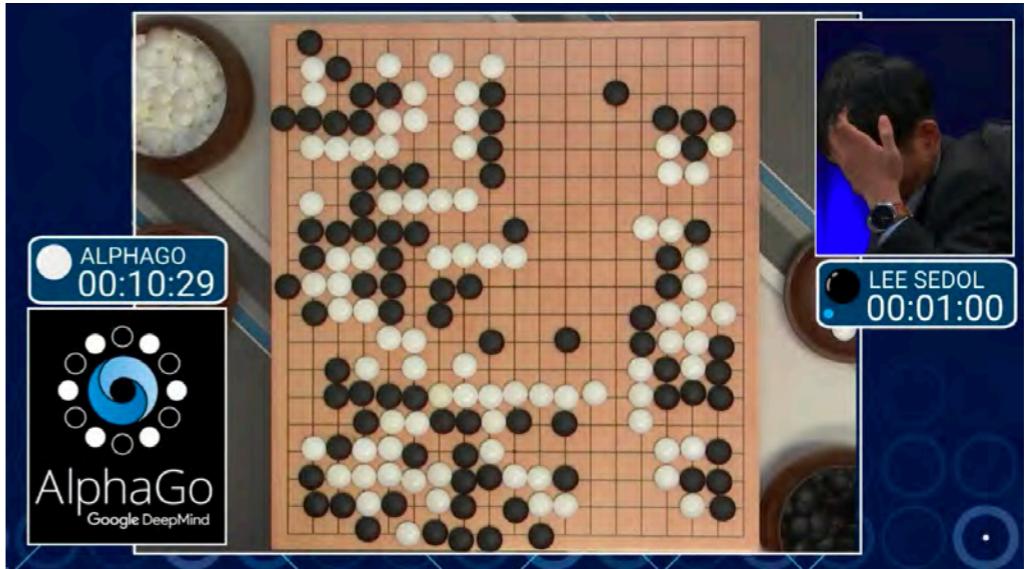
**AlphaFold2:** A step towards tackling the protein folding problem by training semi-supervised deep learning models on publicly available data consisting of ~170,000 protein structures from the protein data bank together with large databases containing protein sequences of unknown structure.

# From predictions to decisions

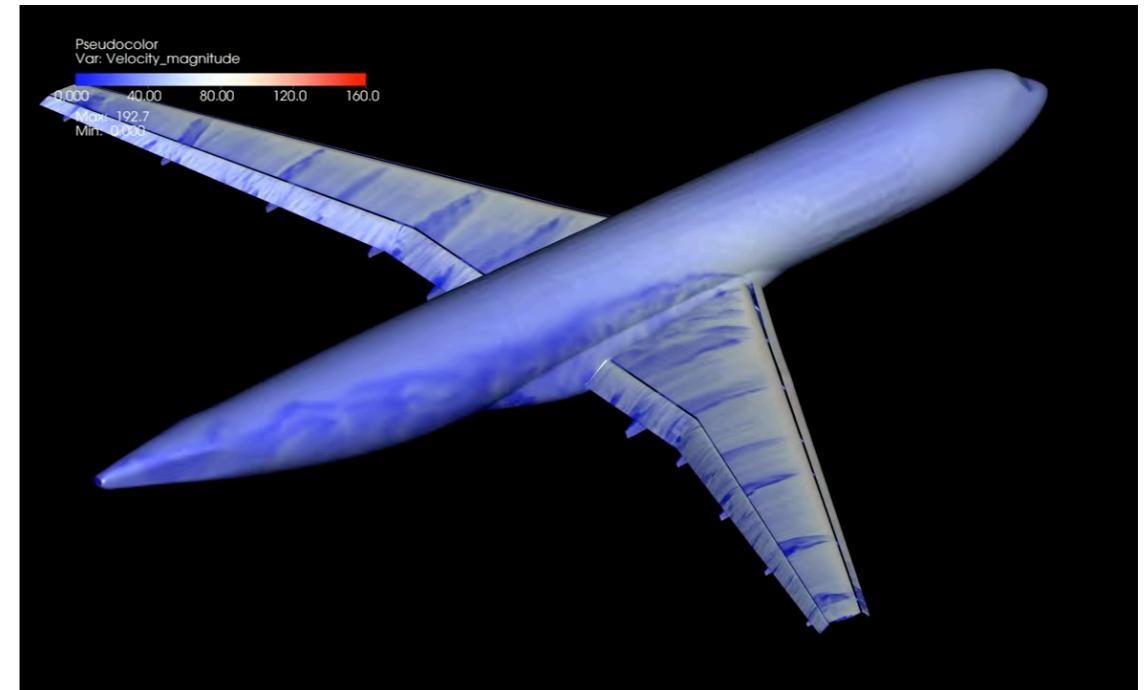
## Google DeepMind's Deep Q-learning

The algorithm will play Atari breakout.

The most important thing to know is that all the agent is given is sensory input (what you see on the screen) and it was ordered to maximize the score on the screen.



# Data- vs Model-driven science

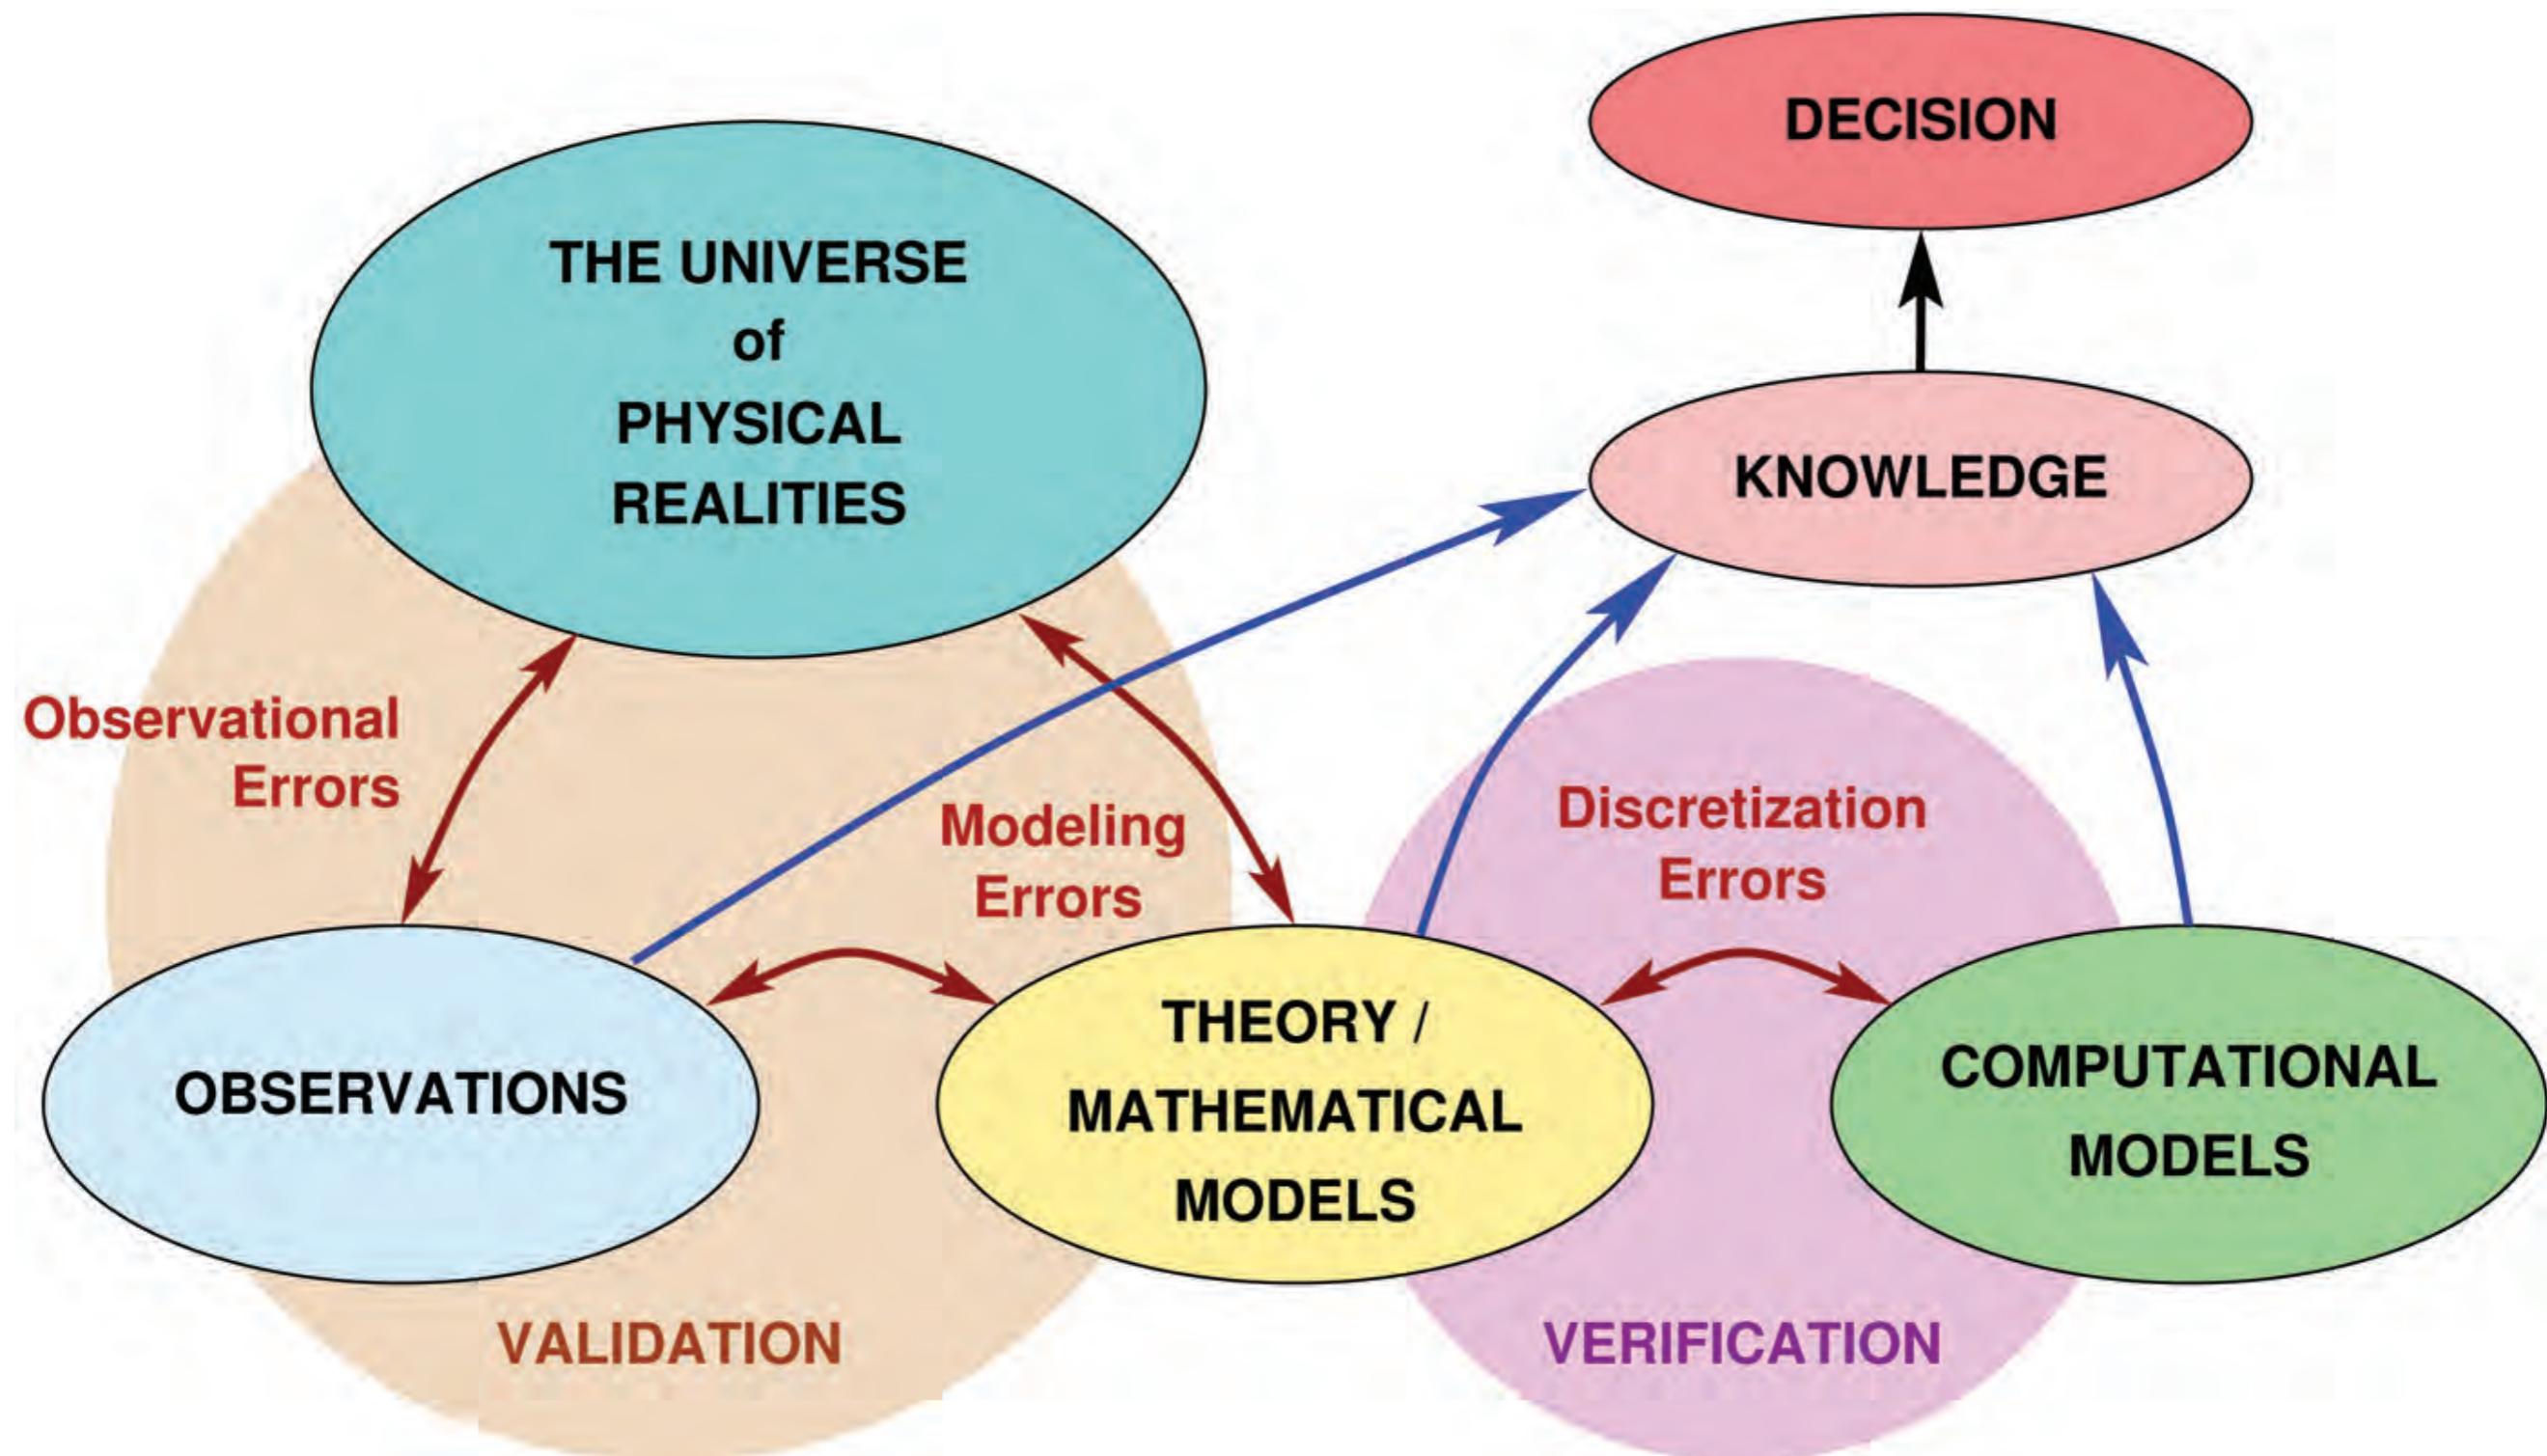


*All models are wrong  
but some are useful*



George E.P. Box

# Predictive science



SPINGER BRIEFS IN STATISTICS

Jordi Vallverdú

Bayesians Versus  
Frequentists  
A Philosophical  
Debate on  
Statistical  
Reasoning

Springer



Aleatoric vs Epistemic Uncertainty

# Course goals

## Foundational Competencies:

- Understand probabilistic modeling and apply Bayesian thinking to engineering problems
- Implement gradient-based optimization algorithms and understand stochastic training methods
- Quantify uncertainty in data-driven models and predictions for engineering applications

## Deep Learning Expertise:

- Design and train neural networks (MLPs) for regression and classification tasks in science and engineering
- Apply convolutional neural networks (CNNs) to spatial data and image analysis
- Understand and implement Transformer architectures and attention mechanisms for sequence data
- Build recurrent neural networks (RNNs, LSTMs) for time-series and dynamical systems

## Physics-Informed AI:

- Understand and implement physics-informed neural networks (PINNs) that incorporate physical constraints
- Apply neural operators to solve engineering problems governed by differential equations
- Integrate domain knowledge and conservation laws into machine learning models

## Generative and Unsupervised Learning:

- Understand generative modeling approaches for data generation and compression
- Apply unsupervised learning techniques to discover patterns in engineering data

## Practical Engineering Skills:

- Combine multiple ML techniques to solve complex engineering problems
- Design and execute computational experiments with proper validation strategies
- Evaluate trade-offs between traditional computational methods and AI approaches
- Implement end-to-end ML pipelines using modern Python frameworks
- Communicate technical results clearly to both technical and non-technical audiences

# Disclaimers & FAQs

## *Is this the right course for me?*

Yes, if you:

- (a) Fulfill the pre-requisites (linear algebra, probability & statistics, optimization, Python programming).
- (b) Want to build foundational knowledge in ML for engineering research and applications.
- (c) Want to sharpen your implementation skills.

No, if you:

- (a) Do not fulfill the pre-requisites.
- (b) Are solely interested in theory.
- (c) Are solely interested in CS applications (e.g. computer vision, NLP).
- (d) Are solely interested in specialized topics (e.g. RL, geometric deep learning).

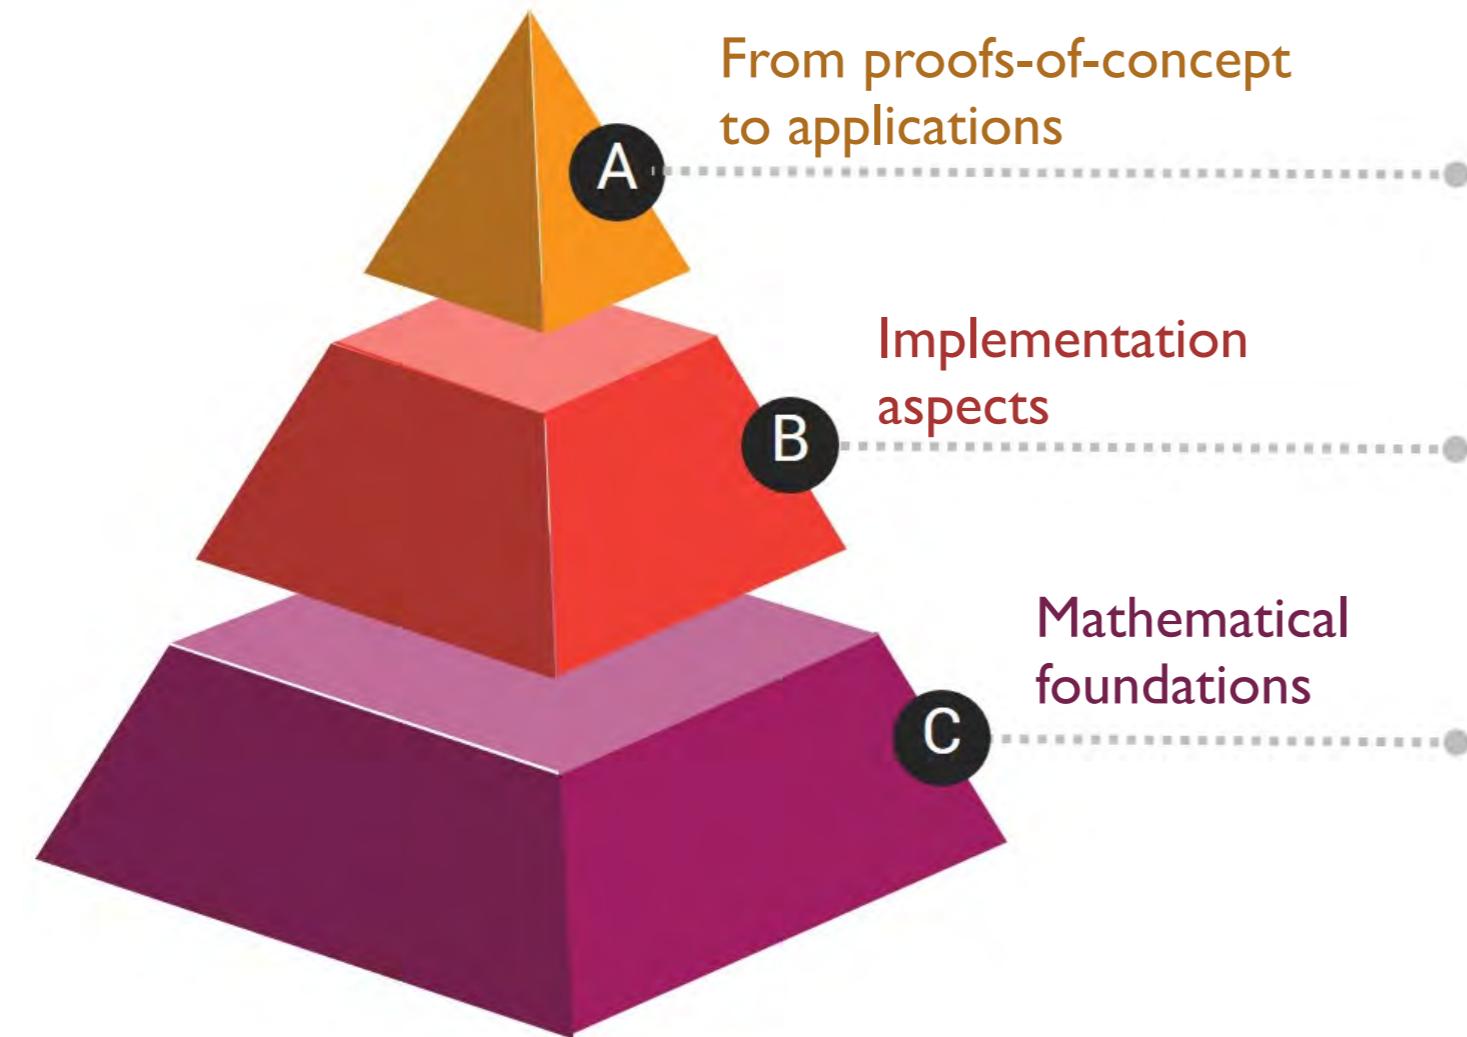
## *I am an experimentalist, what should I expect to learn?*

- (a) How to build predictive models from observational data.
- (b) How to judiciously design and automate experiments.
- (c) How to distill governing laws from experimental data.
- (d) How to deal with aleatoric uncertainty

## *I am a computational scientist, what should I expect to learn?*

- (a) How to accelerate the prediction of large-scale and costly simulations.
- (b) How to design and automate experiments.
- (c) How to calibrate large-scale and costly computational models.
- (b) How to deal with epistemic uncertainty.

# Course philosophy



80% of the homework assignments will be computational primarily based on proof-of-concept studies. Final projects will be geared towards applications.

20% of our lecture time will be devoted on hands-on programming tutorials.

80% of our lecture time will be devoted on discussing the mathematical foundations of modern ML methods.

## Assignments Grade Overall

Your grade is based on your performance during the semester. It includes readings, participation and hands-on tutorials in class (that means you must attend all classes), as well as a final project assignments, broken down as follows.

Assignment	Percentage Value	
Homework	40%	computational proof-of-concept studies
Midterm exam	20%	mathematical foundations
Final project	40%	applications
TOTAL	100%	

$$f:\mathcal{X}\rightarrow\mathcal{Y}$$

# Supervised learning

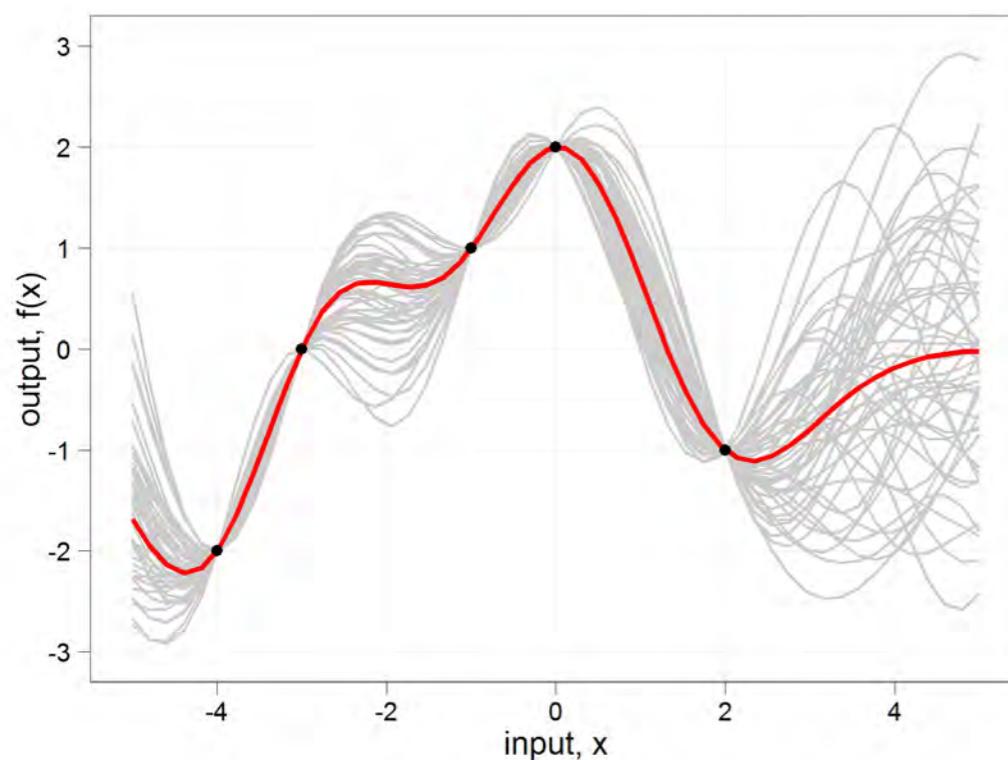
$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

$$\mathcal{D} = \{\mathbf{x}, \mathbf{y}\}, \mathbf{x} \in \mathcal{X}, \mathbf{y} \in \mathcal{Y}$$

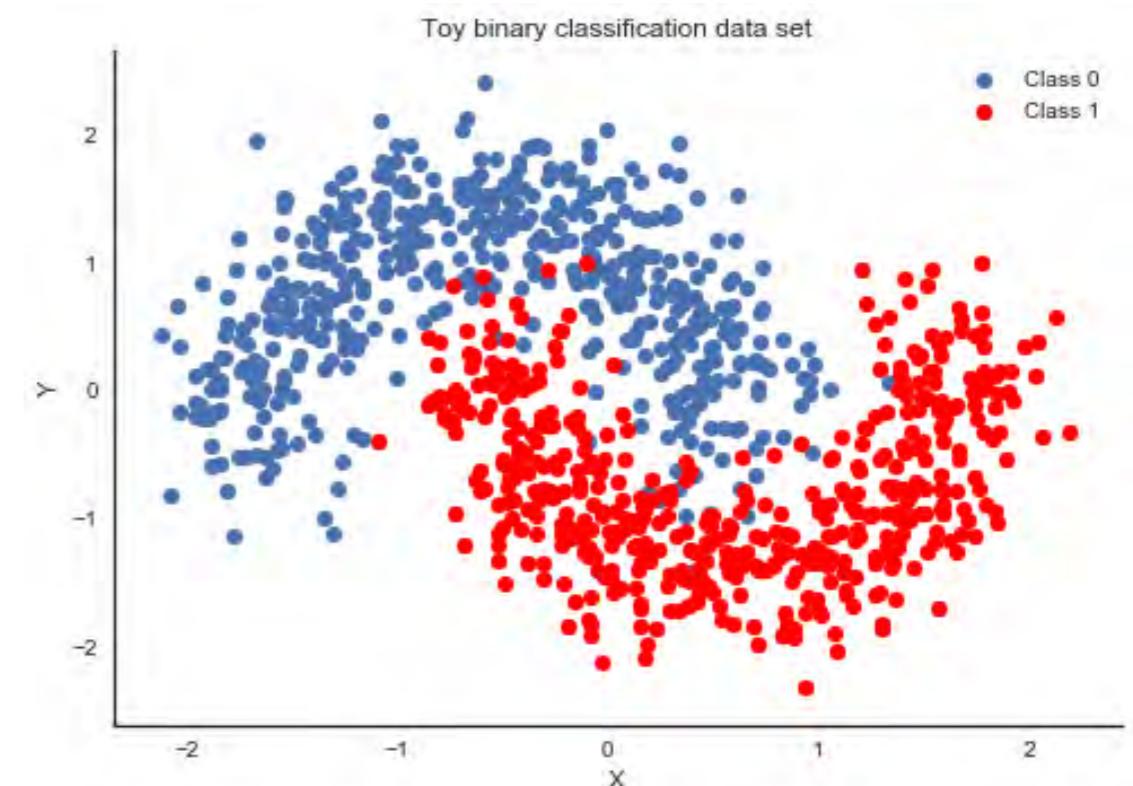
$$\mathbf{y} = f(\mathbf{x}) + \epsilon$$

$$p(f(\mathbf{x}^*) | \mathbf{x}^*, \mathcal{D})$$

Regression

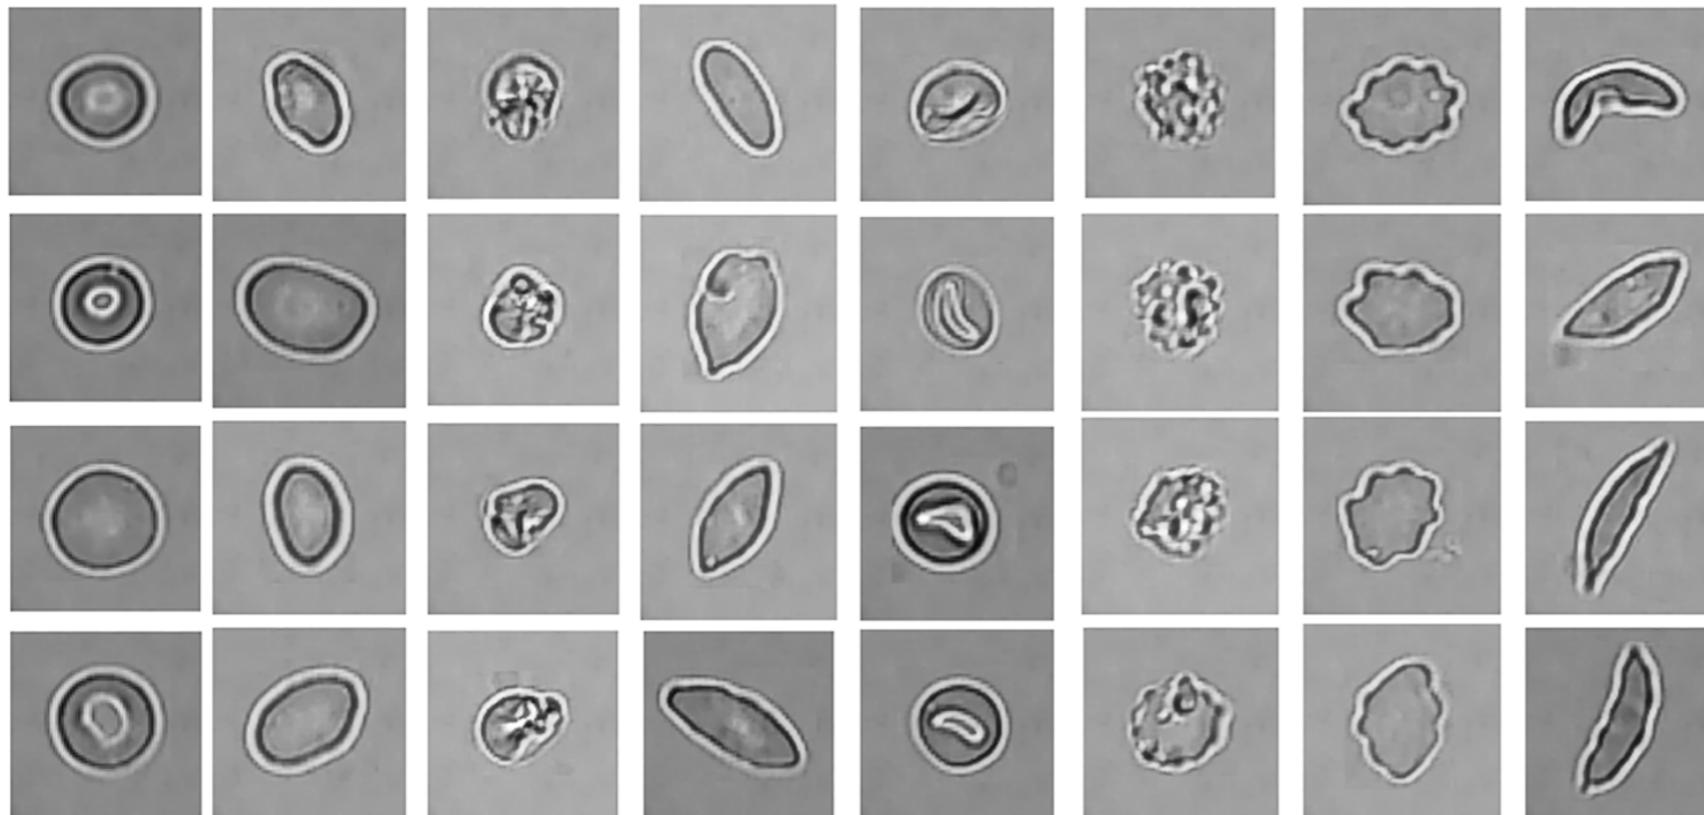


Classification

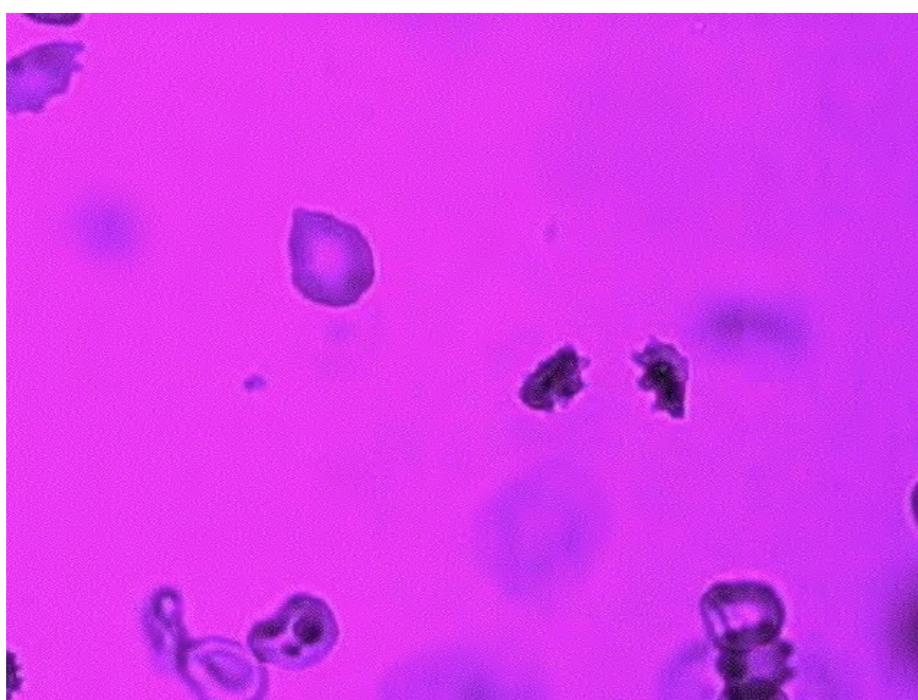


# Classification

Discocytes   Oval   Reticulocytes   Elongated   Stomatocyte   Echinocytes   Granular   Sickle



$$\mathcal{D} = \{x, y\}, \quad x \in \mathcal{X}, \quad y \in \mathcal{Y}$$



$$y = f(x) + \epsilon$$

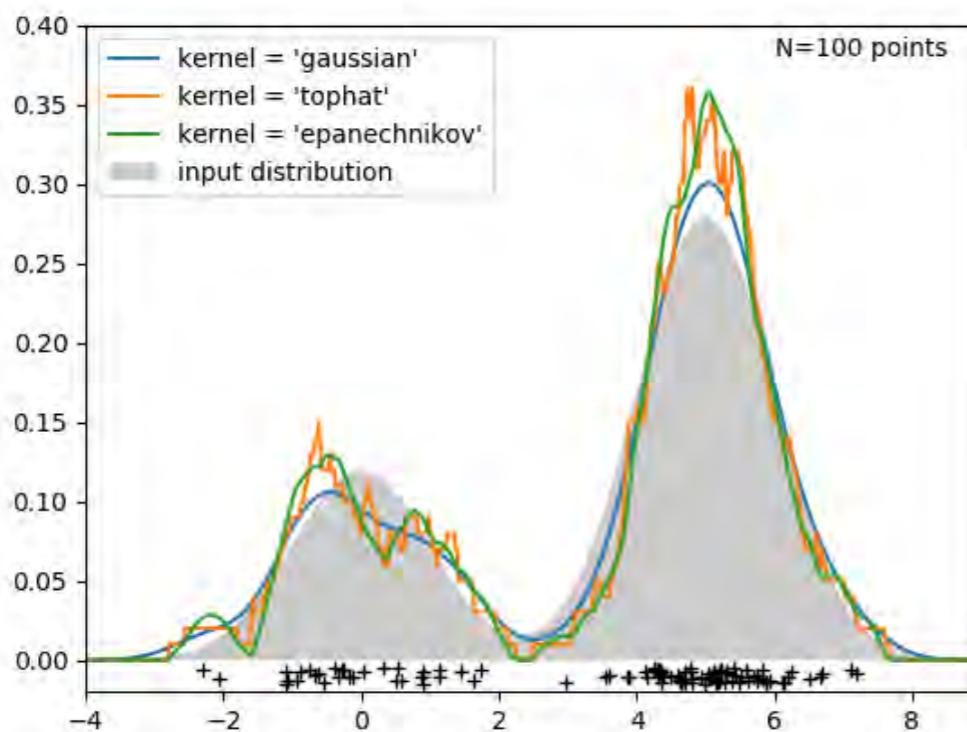
# Unsupervised learning

$$\{y\}, \quad y \in \mathcal{Y}$$

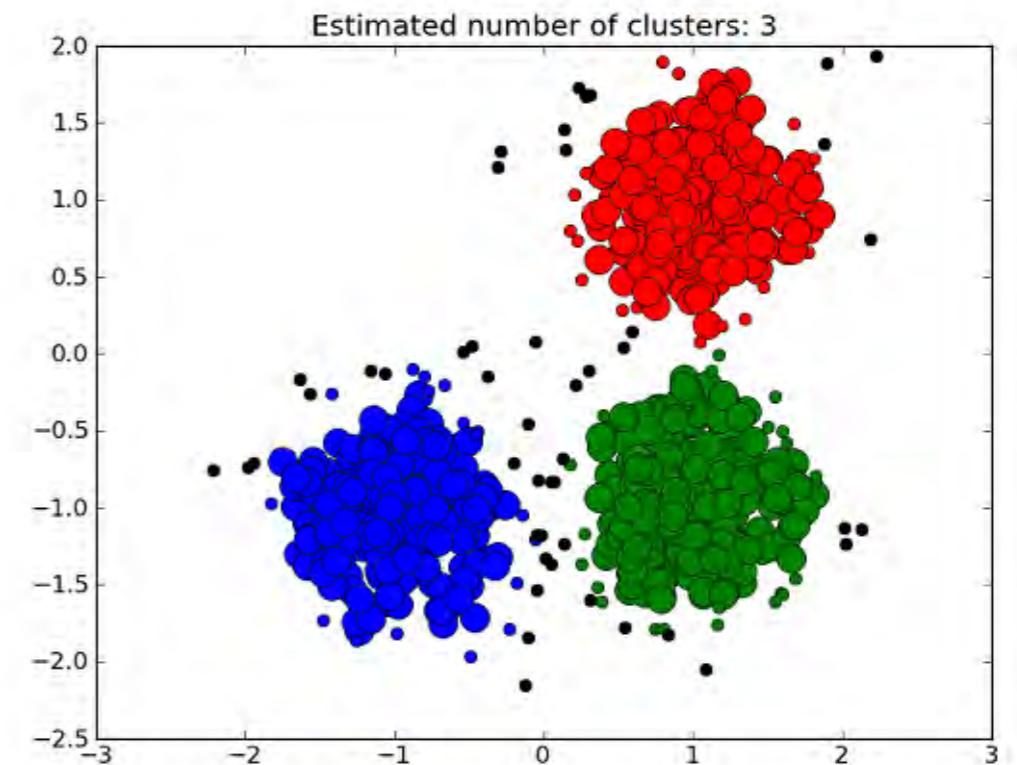
$$y = f(z) + \epsilon$$

$$p(y), \quad z \sim p(z)$$

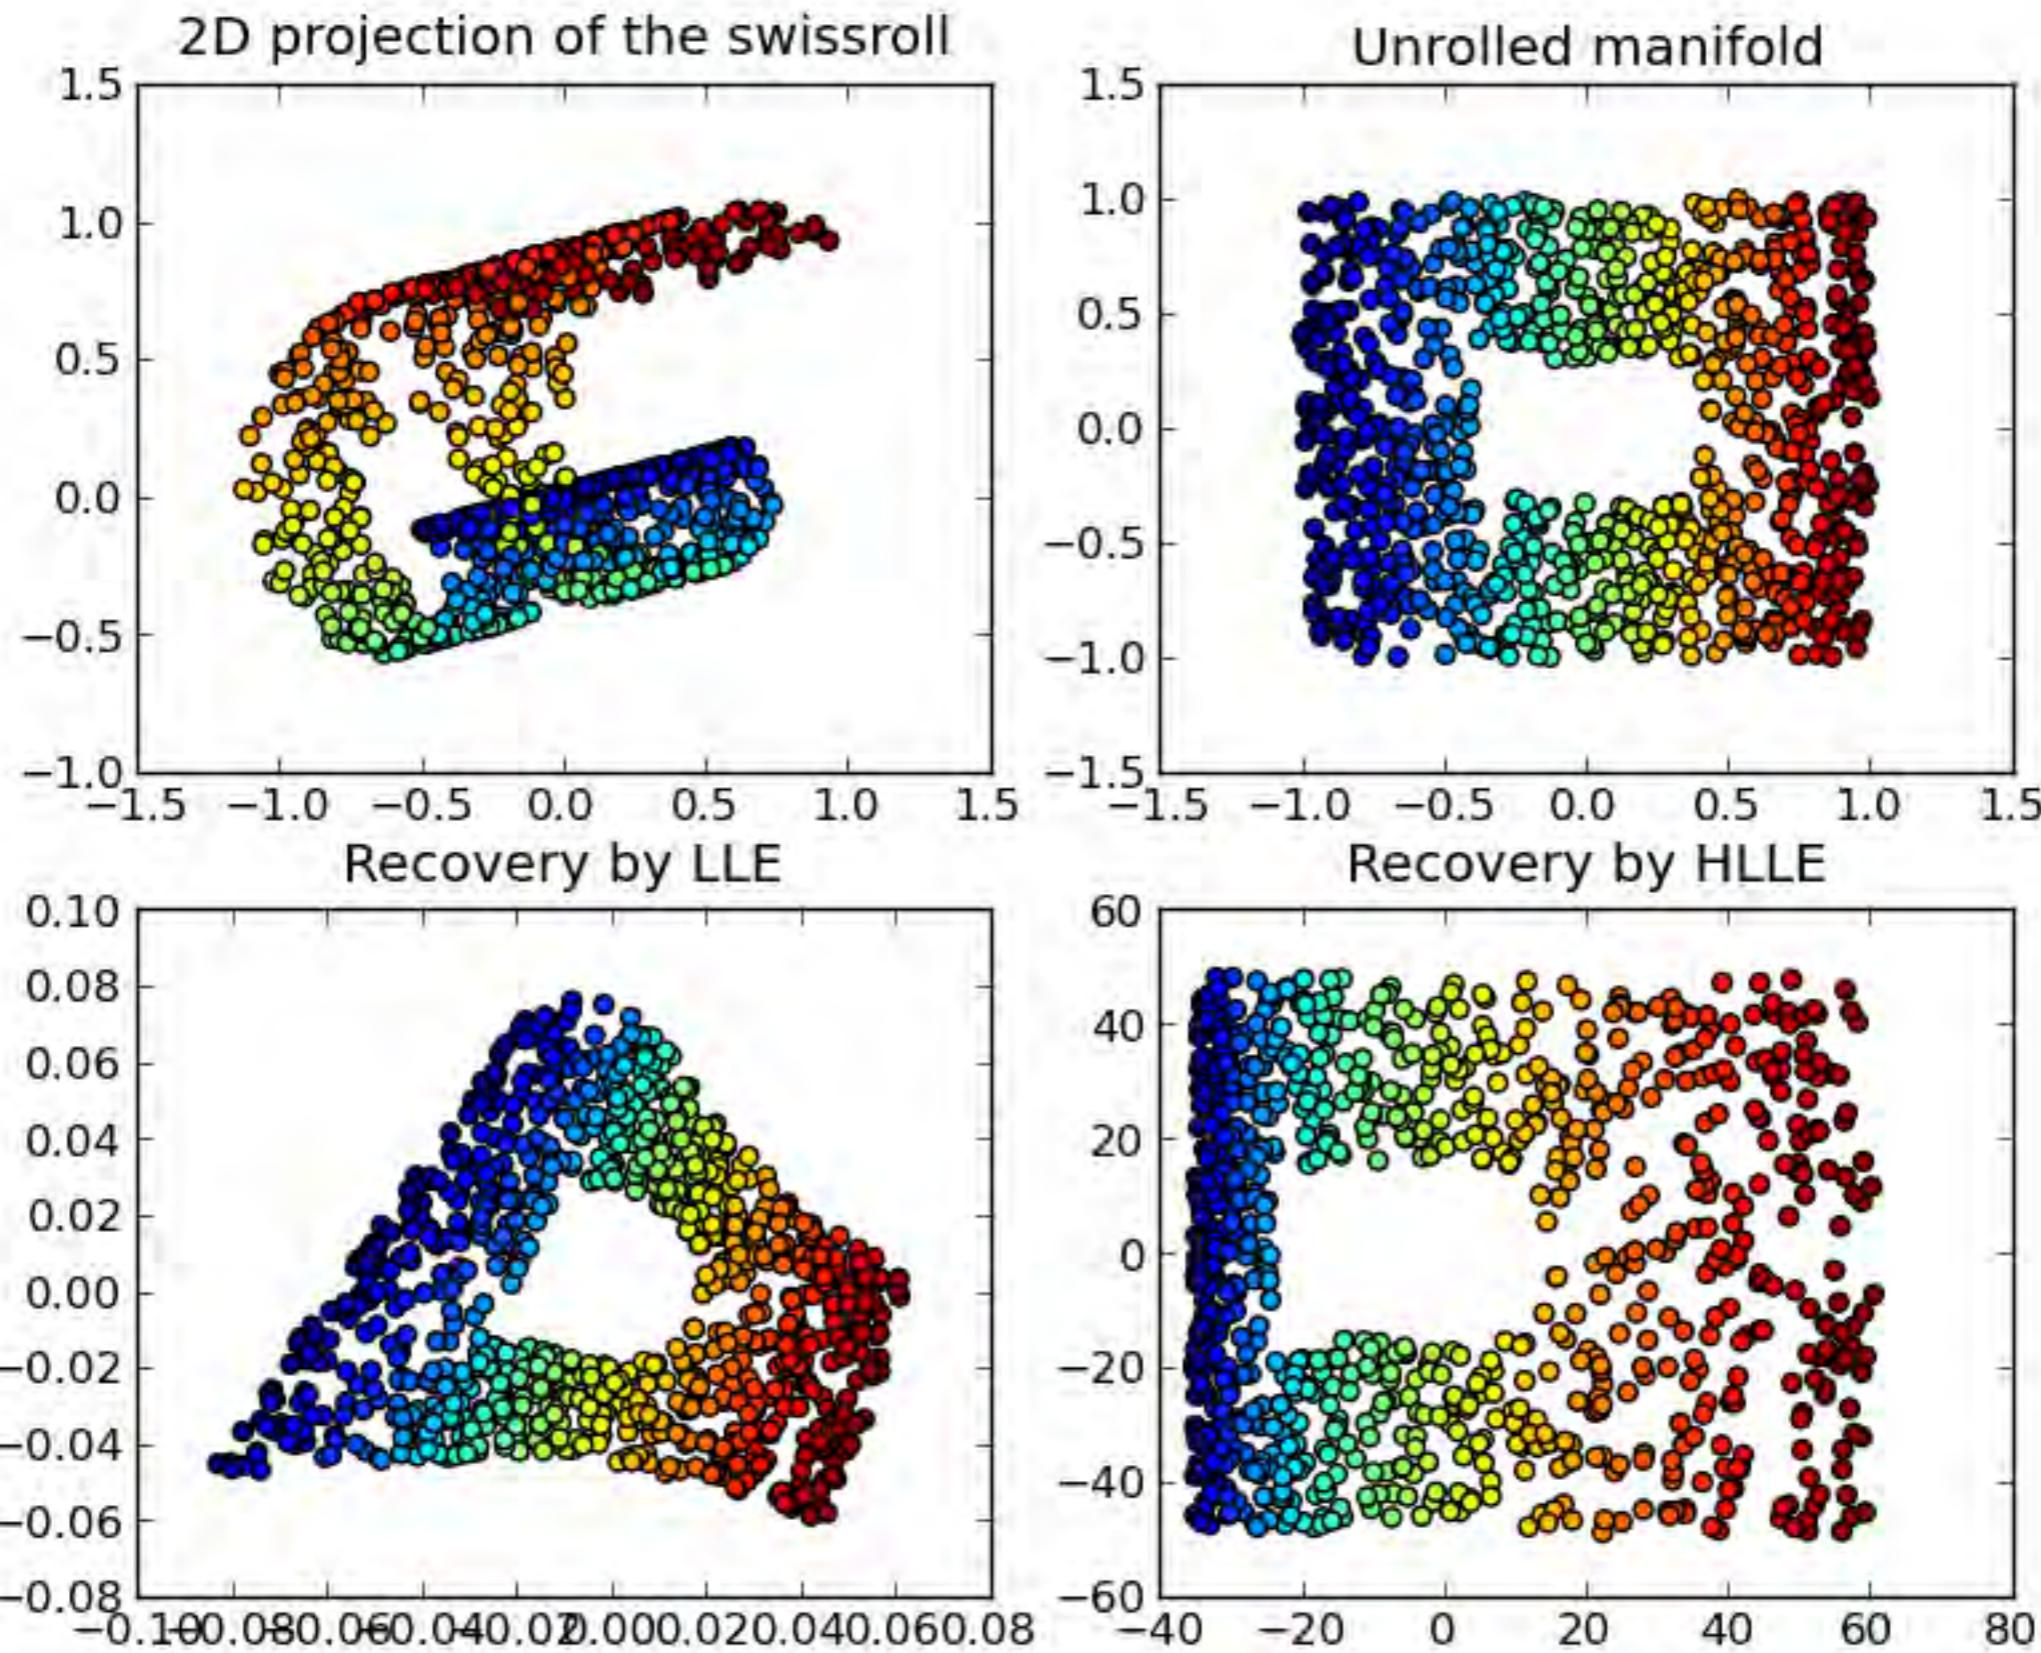
Density estimation



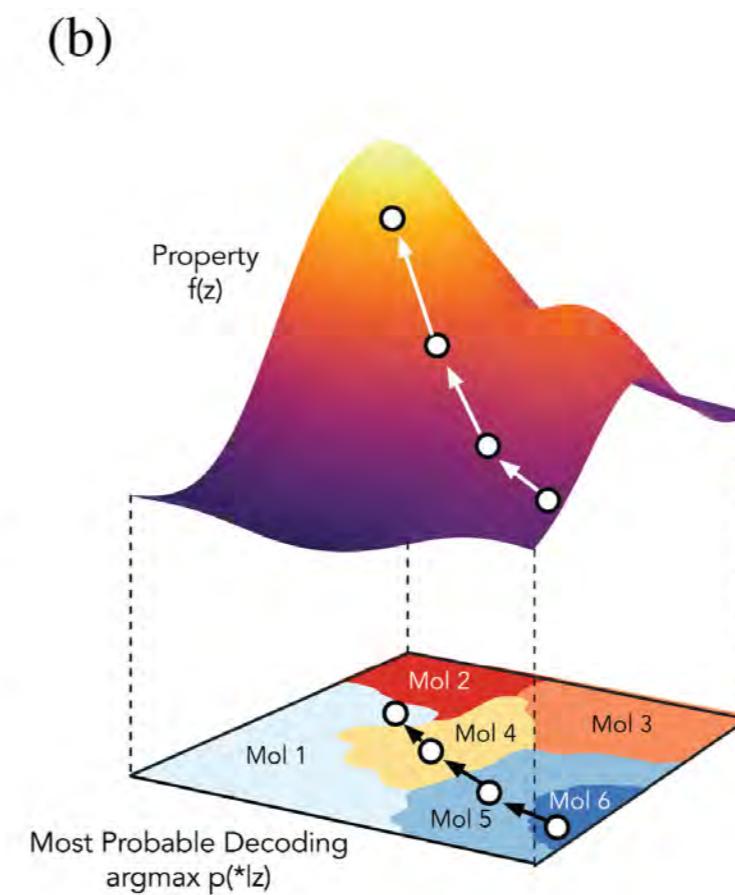
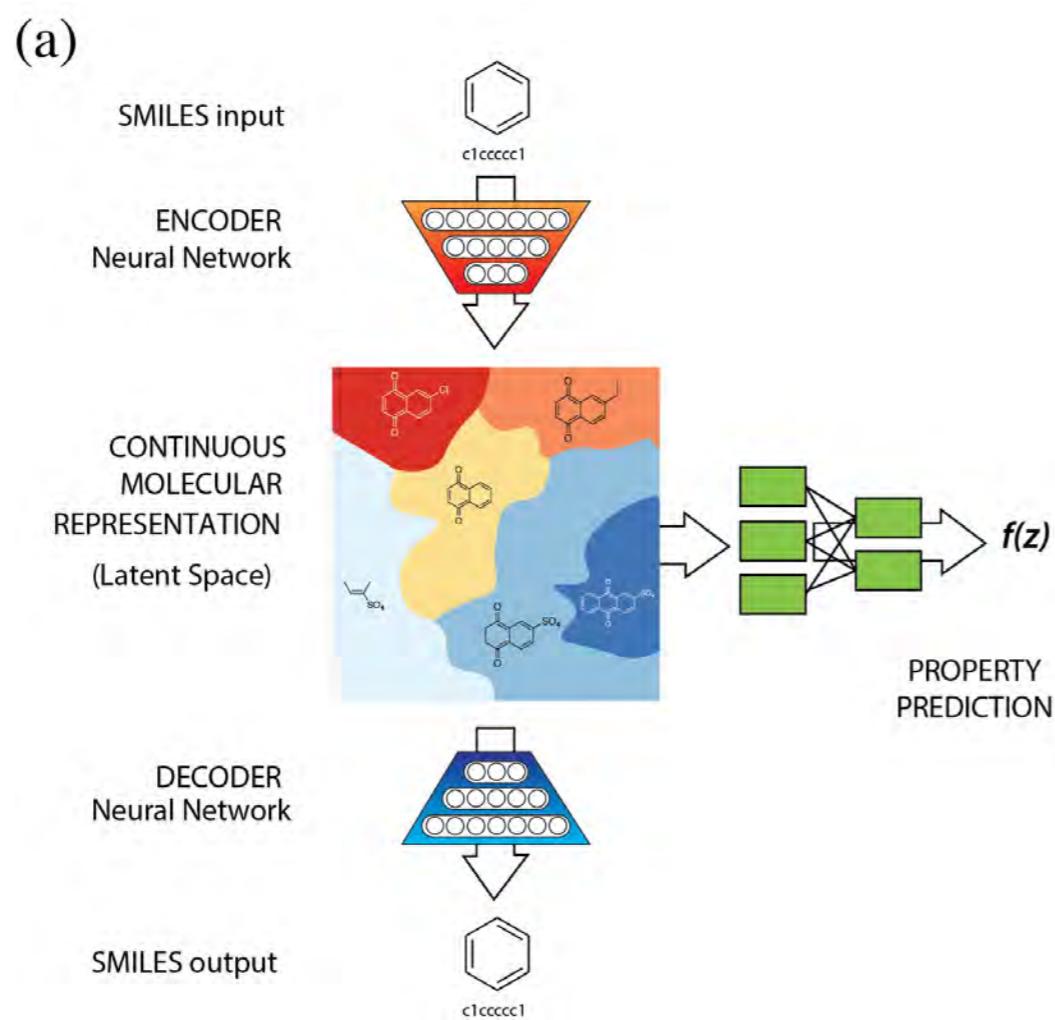
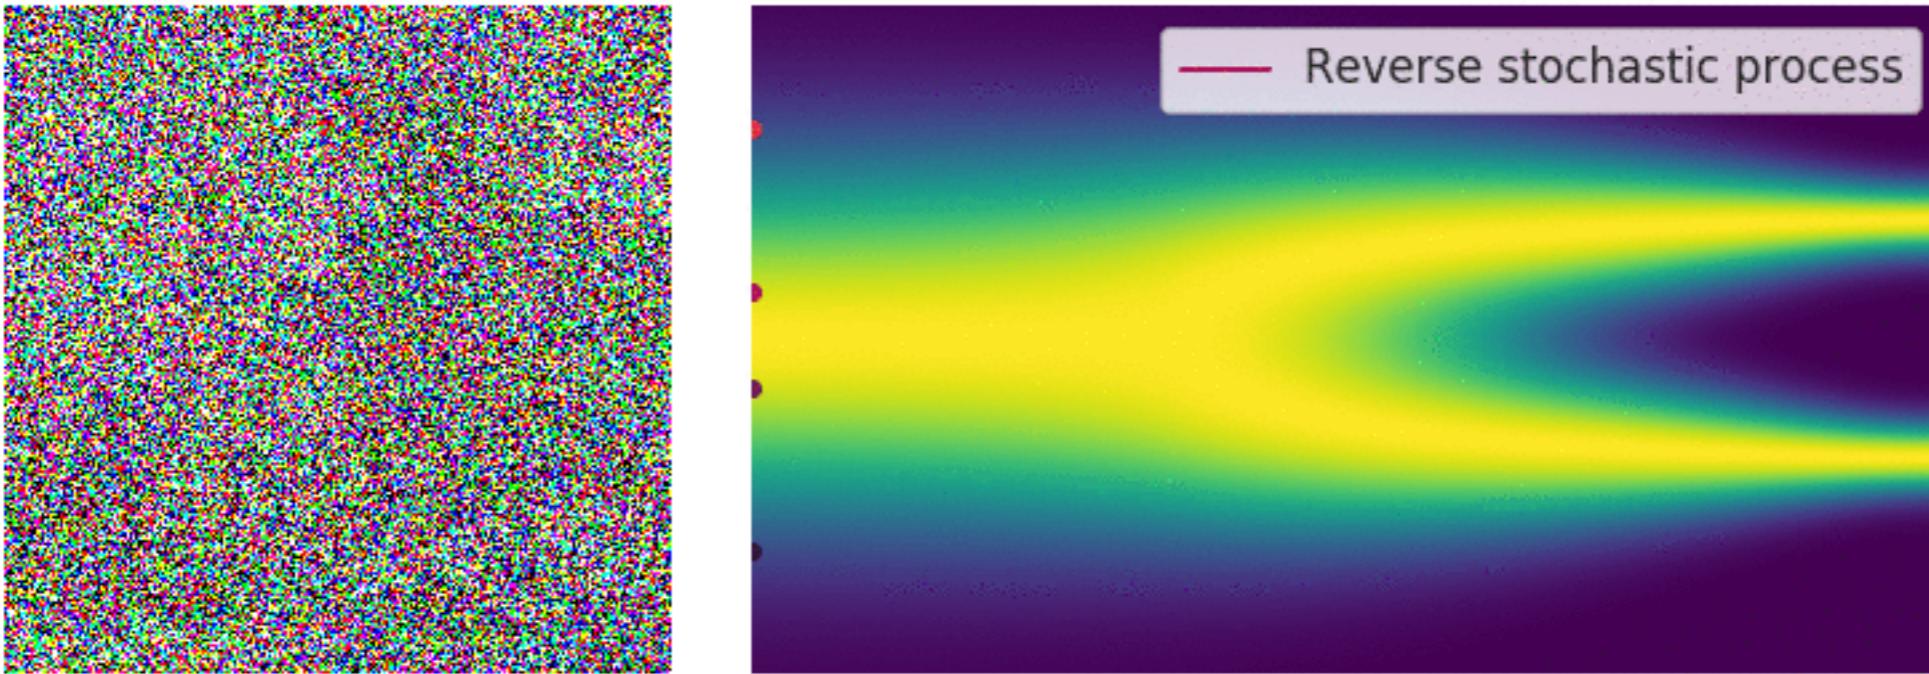
Clustering



# Dimensionality reduction



# Generative Models



## Course logistics

- Duration: 15 Weeks
- Time: Tuesdays and Thursdays, 1:45pm to 3:15pm (Towne 319)
- Dates: January 14 to April 29, 2026
- Instructor office hours: Tuesdays & Thursdays 3.30-4.30pm
- Weekly/bi-weekly HW assignments, mid-term exam, final project.

For more details visit the course website:

<https://www.seas.upenn.edu/~meam4600/>

Slides and code can be found on Github:

<https://github.com/PredictiveIntelligenceLab/MEAM4600>

## **Part I: Foundations**

# Tentative Syllabus

**Week 1:** Introduction to AI in Science and Engineering

**Week 2:** Fundamentals of Machine Learning for Scientific Computing

**Week 3:** Primer on Probability and Statistical Estimation

**Week 4:** Linear & Logistic Regression

## **Part II: Probabilistic Methods**

**Week 5:** Variational Inference

**Week 6:** Monte Carlo Sampling

## **Week 7: Midterm Exam**

## **Part III: Deep Learning Architectures**

**Week 8:** Deep Learning Fundamentals: MLPs and CNNs

**Week 9:** Advanced Architectures: RNNs, GNNs, and Transformers

## **Part IV: Scientific Machine Learning**

**Week 10:** Kernel Methods & Gaussian Process Regression

**Week 11:** Physics-Informed Neural Networks & Neural Operators

**Week 12:** Bayesian Optimization and Active Learning

## **Part V: Unsupervised and Generative Methods**

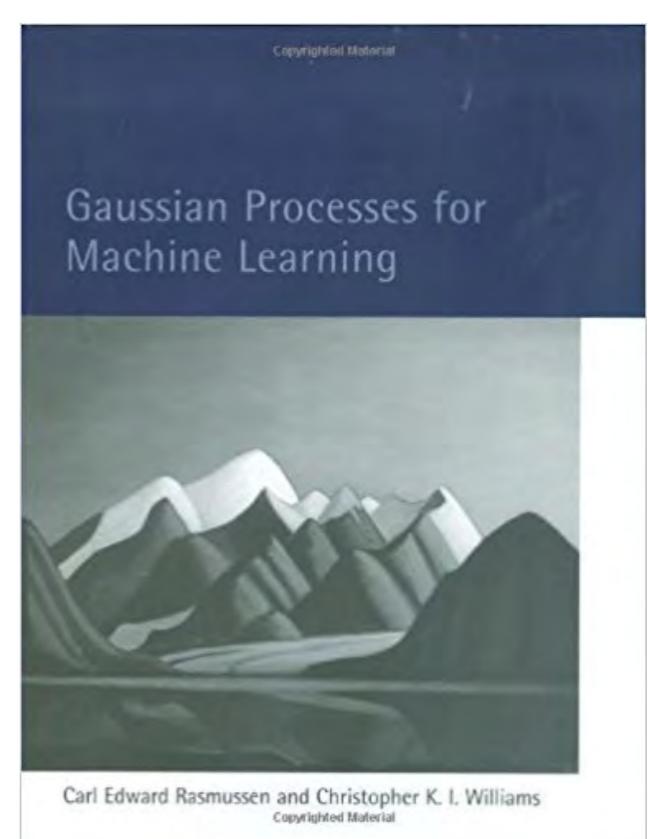
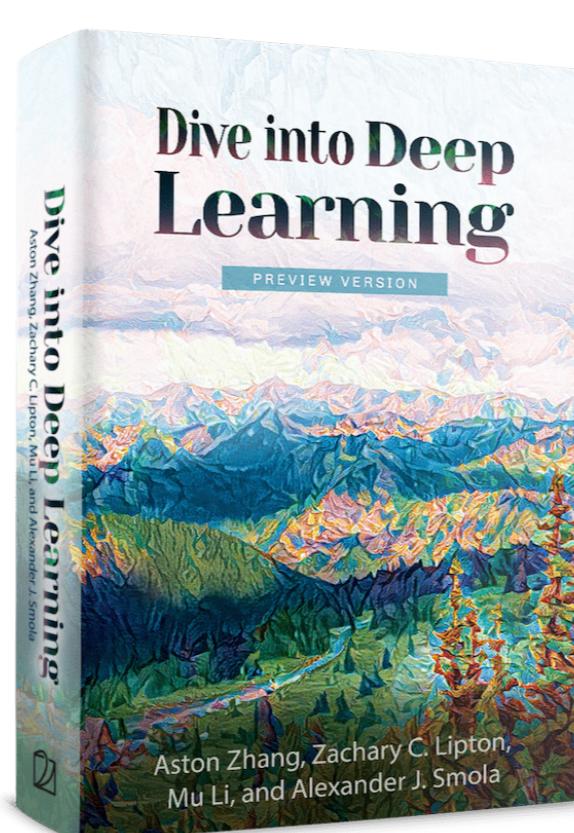
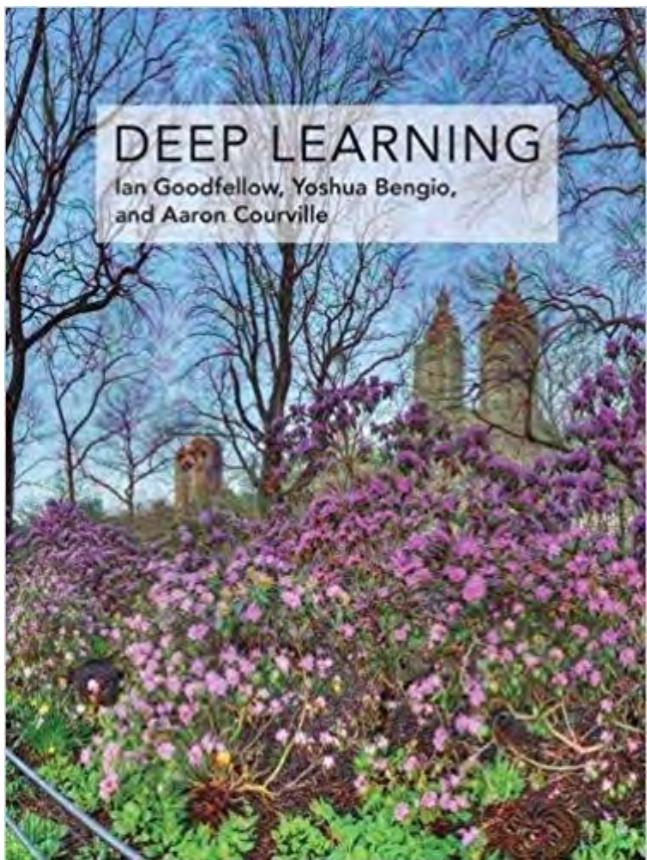
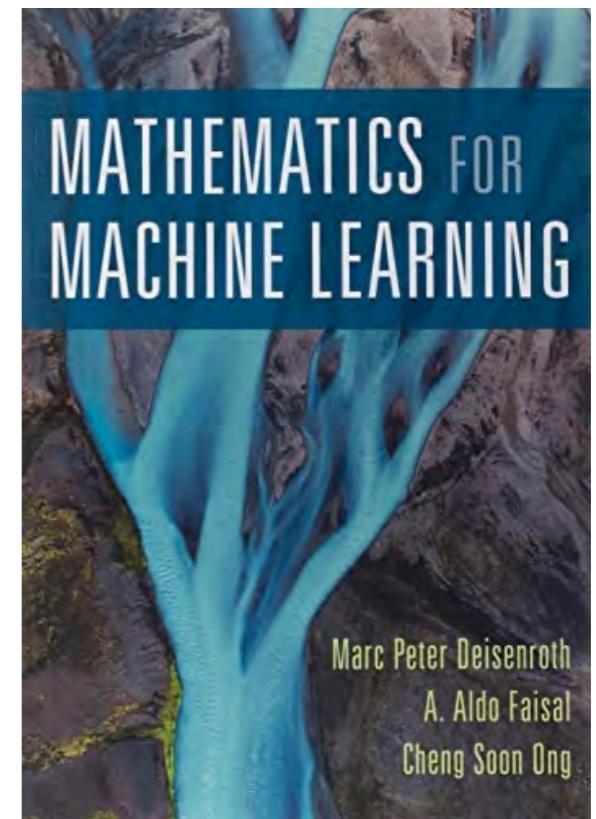
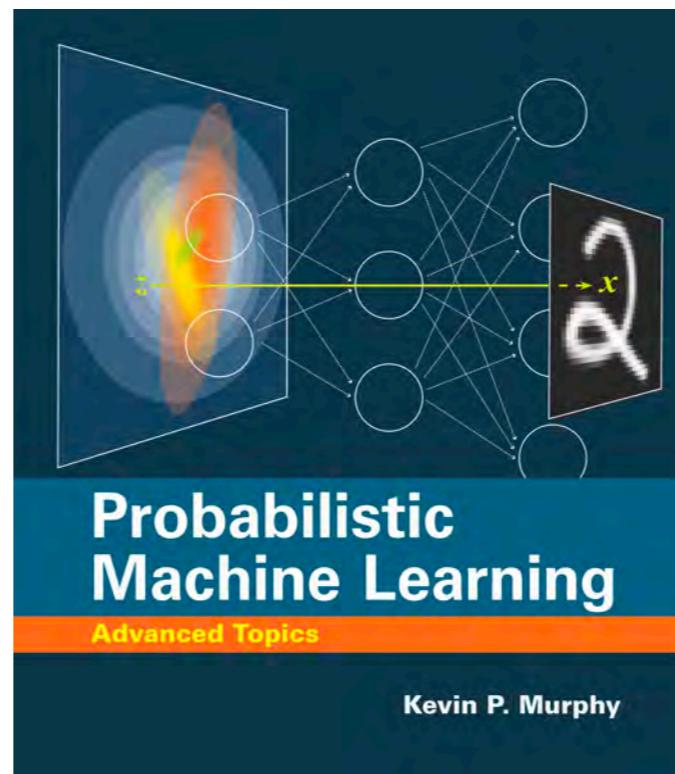
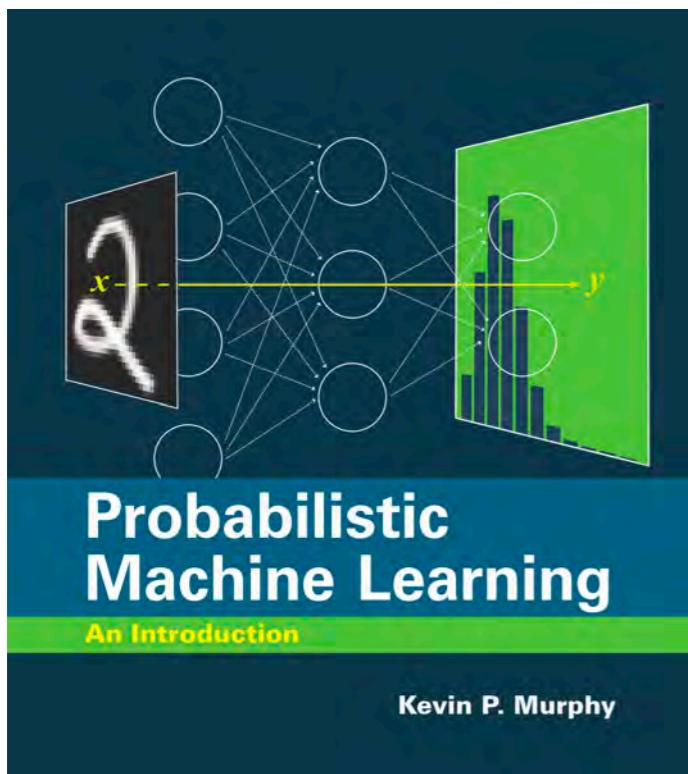
**Week 13:** Unsupervised Learning: PCA and VAEs

**Week 14:** Generative Models

## **Week 15: Final Project Presentations**

**Note:** Each week includes lecture content, hands-on coding exercises, and real-world engineering examples. Homework assignments are distributed throughout the semester to reinforce concepts progressively.

# Textbooks



# Reading

- Introduction
- Syllabus review
- Presentation of diverse applications that showcase the use of data, modeling, and scientific computation.
- Overview of the main themes and goals of this course
- Primer on Probability and Statistics

## Reading

- The Emergence of Predictive Computational Science:
    - Computer predictions with quantified uncertainty ([Oden, Moser, & Ghattas, 2010](#))
    - [Lecture](#) by J.T Oden.
  - Review papers on recent advances in machine learning:
    - Probabilistic machine learning and artificial intelligence ([Ghahramani, 2015](#))
    - Deep learning ([LeCun, Bengio, & Hinton, 2015](#))
    - Machine learning: Trends, perspectives, and prospects ([Jordan & Mitchell, 2015](#))
  - Scientific computing in Python:
    - [Lectures and code](#) by Robert Johansson.
1. Oden, T., Moser, R., & Ghattas, O. (2010). Computer predictions with quantified uncertainty, part I. *SIAM News*, 43(9), 1–3.
  2. Ghahramani, Z. (2015). Probabilistic machine learning and artificial intelligence. *Nature*, 521(7553), 452–459.
  3. LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436–444.
  4. Jordan, M. I., & Mitchell, T. M. (2015). Machine learning: Trends, perspectives, and prospects. *Science*, 349(6245), 255–260.

# VSCode

The screenshot shows the Visual Studio Code interface with the following details:

- File Explorer:** Shows the project structure for "CREATE-REACT-APP". The "serviceWorker.js" file is selected in the "src" folder.
- Code Editor:** Displays the "serviceWorker.js" code. A tooltip is open over the "window.a('load', () => {" line, listing various global objects like addEventListener, alert, applicationCache, etc.
- Terminal:** Shows the output of a build command:

```
Compiled successfully!
```

You can now view **create-react-app** in the browser.

Local: http://localhost:3000/  
On Your Network: http://192.168.86.138:3000/

Note that the development build is not optimized.  
To create a production build, use **yarn build**.
- Status Bar:** Shows the current branch as "master\*", file count as "0", and other status indicators.

# Jupyter notebook

The screenshot shows a Jupyter Notebook interface with the following components:

- Header:** The title "jupyter Untitled (unsaved changes)" is at the top left, and a Python 2 logo is at the top right.
- Toolbar:** A horizontal toolbar with icons for file operations (New, Open, Save, etc.), cell selection (Cell, Kernel), and help.
- Text Cell:** A text cell containing the heading "Simple Jupyter demo" and a descriptive paragraph about markdown rendering.
- Code Cell:** An execution cell labeled "In [57]:" containing Python code to generate three random numbers. The output shows the results: 0.10564822904, 0.153941700348, and 0.518503128416.
- Text Cell:** Another text cell containing the text "Here is another text cell, with some *formatting*.  
This cell has text formatted using the markdown language, which gets rendered like regular html."

# Google colab

The screenshot shows a Google Colab notebook titled "face\_detection.ipynb". The notebook interface includes a top bar with tabs for "face\_detection - Google Drive" and "face\_detection.ipynb - Colaboratory". The main area displays a list of steps for running machine learning projects:

- Using Google Colab to run our machine learning projects.
- Change the directory to face\_detection folder  
First create a folder and name it as `face_detection`

```
from google.colab import drive
drive.mount('/content/drive/')
```
- use `os.chdir` to connect to the face\_detection directory  

```
from os import chdir
chdir('/content/drive/My Drive/face_detection')
```
- Clone your project in the google drive  

```
!git clone https://github.com/masouduut94/Pytorch_Retinaface.git
```

Cloning into 'Pytorch\_Retinaface'...  
remote: Enumerating objects: 26, done.  
remote: Counting objects: 100% (26/26), done.  
remote: Compressing objects: 100% (18/18), done.  
remote: Total 146 (delta 7), reused 23 (delta 5), pack-reused 120  
Receiving objects: 100% (146/146), 12.34 MiB | 24.73 MiB/s, done.  
Resolving deltas: 100% (46/46), done.
- Doing a little bit cleaning.  

```
!mv Pytorch_Retinaface/* ./
!rm -r Pytorch_Retinaface/
```
- Get the pretrained weight files  
1 - We open the link to the pretrained models.  
2 - We right click on the weight file and choose "Add Shortcut to Drive"  
3 - Then we choose the `/content/drive/My Drive/face_detection` directory

Disk: 37.15 GB available

<https://colab.research.google.com/notebooks/intro.ipynb>